



# Regularization for Deep Learning

---

ADVANCED ARTIFICIAL INTELLIGENCE  
JUCHEOL MOON

---

1

## Regularization

---

- Regularization is any modification we make to a learning algorithm that is intended to reduce
  - its generalization(test) error
  - but not its training error
- Regularization strategies
  - put extra constraints on a machine learning model
  - add extra terms in the objective function

2

2

# Parameter Norm Penalties

---

- Limiting the capacity of models by adding a parameter norm penalty  $\Omega(\theta)$  to the objective function  $J$ .

$$\tilde{J}(\vec{\theta}; \vec{X}, \vec{y}) = J(\vec{\theta}; \vec{X}, \vec{y}) + \alpha \Omega(\vec{\theta})$$

- where  $\alpha \in [0, \infty)$  is a hyperparameter.
- $\alpha = 0$  results in no regularization,
- Large  $\alpha$  correspond to more regularization.
- It is sometimes desirable to use a separate penalty with a different  $\alpha$  coefficient for each layer of the network.

3

3

## $L^2$ Parameter Regularization

---

- Also known as weight decay and ridge regularization

$$\tilde{J}(\vec{w}; \vec{X}, \vec{y}) = J(\vec{w}; \vec{X}, \vec{y}) + \frac{\alpha}{2} \vec{w}^T \vec{w}$$

$$\nabla_{\vec{w}} \tilde{J}(\vec{w}; \vec{X}, \vec{y}) = \nabla_{\vec{w}} J(\vec{w}; \vec{X}, \vec{y}) + \alpha \vec{w}$$

$$\underline{\vec{w}} \leftarrow \vec{w} - \varepsilon (\alpha \vec{w} + \nabla_{\vec{w}} J(\vec{w}; \vec{X}, \vec{y}))$$

$$\vec{w} \leftarrow \underline{\vec{w}} - \varepsilon \alpha \vec{w} - \varepsilon \nabla_{\vec{w}} J(\vec{w}; \vec{X}, \vec{y}) = \underline{(1 - \varepsilon \alpha) \vec{w}} - \varepsilon \nabla_{\vec{w}} J(\vec{w}; \vec{X}, \vec{y})$$

- shrink the weight vector by a constant factor on each step

4

4

# Weight Decay as Constrained Optimization

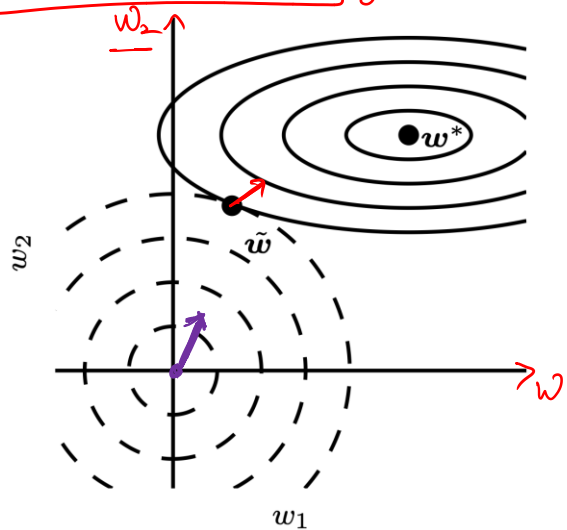
$$\vec{w} \leftarrow \boxed{(1 - \epsilon\alpha)\vec{w}} - \epsilon \nabla_{\vec{w}} J(\vec{w}; \vec{X}, \vec{y}) \quad \vec{g}$$

$$0 < \epsilon \ll 1$$

$$0 < \alpha \ll 1$$

$$0 < 2\alpha \ll 1$$

$$\underline{1 - \epsilon\alpha < 1}$$



5

5

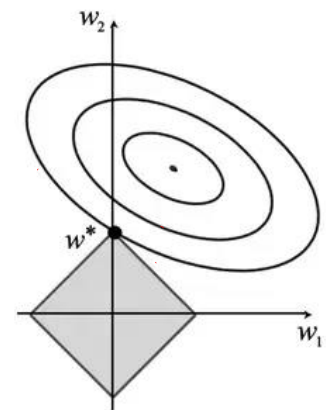
# $L^1$ Parameter Regularization

▪ Also known as LASSO regularization

$$\tilde{J}(\vec{w}; \vec{X}, \vec{y}) = J(\vec{w}; \vec{X}, \vec{y}) + \alpha \|\vec{w}\|_1$$

$$\nabla_{\vec{w}} \tilde{J}(\vec{w}; \vec{X}, \vec{y}) = \nabla_{\vec{w}} J(\vec{w}; \vec{X}, \vec{y}) + \alpha \text{Sign}(\vec{w})$$

$$\vec{w} \leftarrow \vec{w} - \epsilon \alpha \text{sign}(\vec{w}) - \epsilon \nabla_{\vec{w}} J(\vec{w}; \vec{X}, \vec{y})$$



6

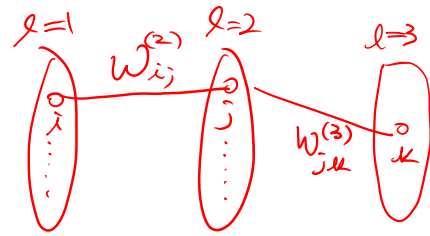
6

# $L^2$ Regularization in Neural Network

$$\tilde{J}(\vec{W}, \vec{b}) = \underbrace{\frac{1}{m} \sum_i^m L(\hat{y}^{(i)}, y^{(i)})}_{\text{J}} + \underbrace{\frac{\alpha}{2} \sum_l^L \|\vec{W}^{(l)}\|_F^2}_{\text{layer } \Omega}$$

▪ Frobenius norm

$$\|\vec{W}^{(l)}\|_F^2 = \sum_i \sum_j (w_{ij}^{(l)})^2$$



7

7

# $L^2$ Regularization in Neural Network

$$\tilde{J}(\vec{W}, \vec{b}) = \frac{1}{m} \sum_i^m L(\hat{y}^{(i)}, y^{(i)}) + \frac{\alpha}{2} \sum_l^L \|\vec{W}^{(l)}\|_F^2$$

▪ In deep neural network

▪  $\alpha \uparrow$ : underfitting,  $\alpha \downarrow$ : overfitting

$$\vec{w} \leftarrow (1 - \epsilon \alpha) \vec{w} - \epsilon D_{\vec{w}} J$$

8

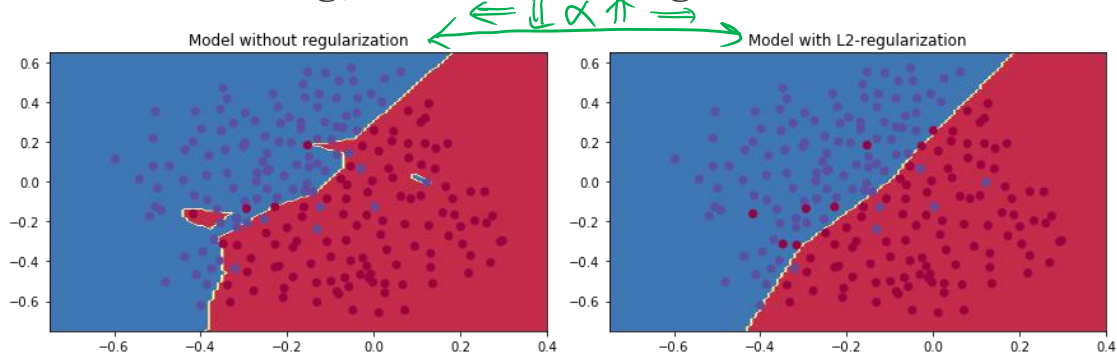
8

# $L^2$ Regularization in Neural Network

$$J(\vec{W}, \vec{b}) = \frac{1}{m} \sum_i^m L(\hat{y}^{(i)}, y^{(i)}) + \frac{\alpha}{2} \sum_l^L \|\vec{W}^{(l)}\|_F^2$$

▪ In deep neural network

▪  $\alpha \uparrow$ : underfitting,  $\alpha \downarrow$ : overfitting

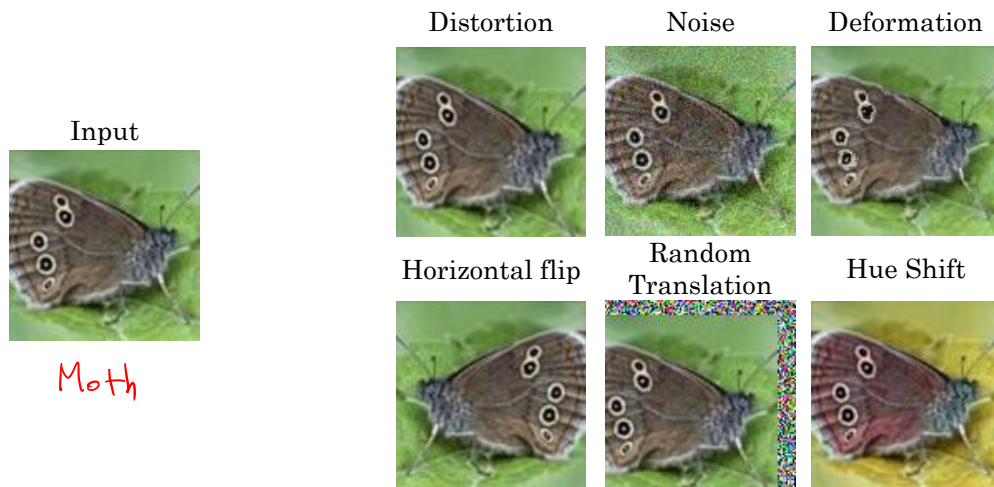


9

9

## Dataset Augmentation

▪ The best way to make a machine learning model generalize better is to train it on more data.

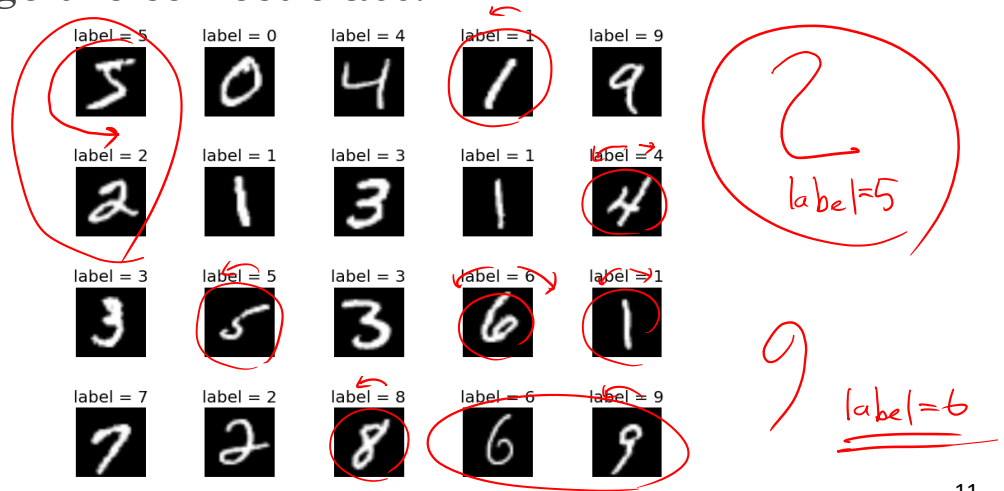


10

10

# Dataset Augmentation

- One must be careful not to apply transformations that would change the correct class.

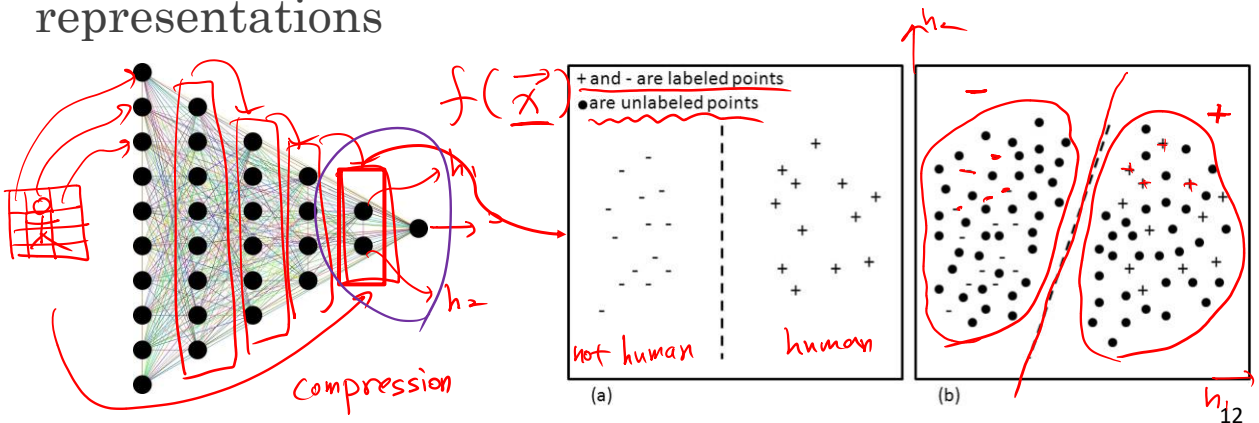


11

11

# Semi-Supervised Learning

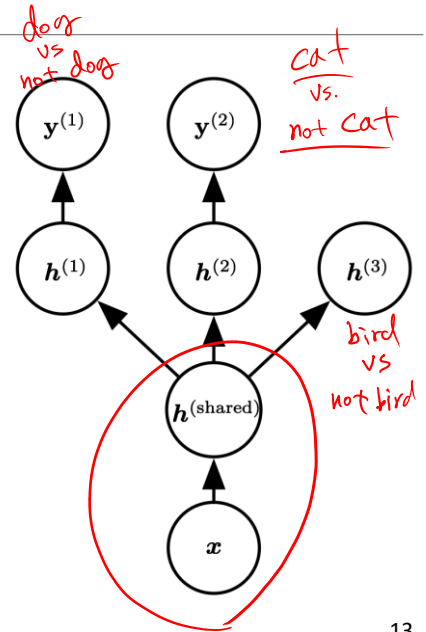
- Semi-supervised learning usually refers to learning a representation  $h = f(x)$ 
  - Examples from the same class have similar representations



12

# Multi-Task Learning

- The model can generally be divided into two kinds of parts and associated parameters:
  - Task-specific parameters (which only benefit from the examples of their task to achieve good generalization)
  - Generic parameters, shared across all the tasks (which benefit from the pooled data of all the tasks).

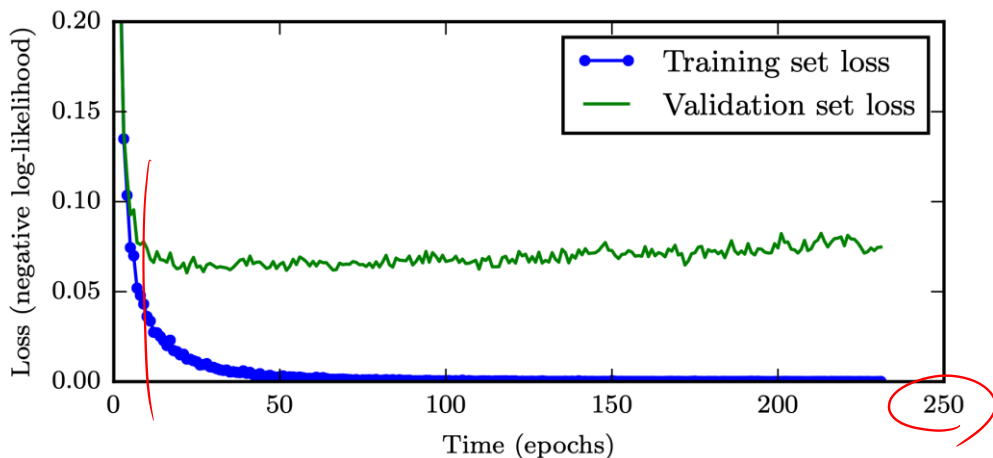


13

13

# Early Stopping

- We often observe that training error decreases steadily over time, but validation set error begins to rise again.



14

14

## Early Stopping version 1

---

- Every time the error on the validation set improves, we store a copy of the model parameters.
- When the training algorithm terminates, we return these parameters, rather than the latest parameters.
- An additional cost to early stopping is the need to maintain a copy of the best parameters.
  - This cost is generally negligible.
- Early stopping requires a validation set, which means some training data is not fed to the model.

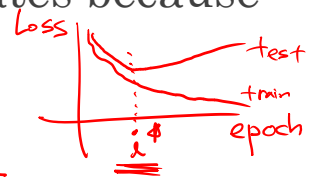
15

15

## Early Stopping version 2

---

- Perform extra training after the initial training with early stopping has completed.)
- In this second training pass, we train for the same number of steps as the early stopping procedure determined was optimal in the first pass.
- On the second round of training, each pass through the dataset will require more parameter updates because the training set is bigger.



16

16



# Early Stopping version 2

---

**Algorithm 7.2** A meta-algorithm for using early stopping to determine how long to train, then retraining on all the data.

---

Let  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  be the training set.

Split  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  into  $(\mathbf{X}^{(\text{subtrain})}, \mathbf{X}^{(\text{valid})})$  and  $(\mathbf{y}^{(\text{subtrain})}, \mathbf{y}^{(\text{valid})})$  respectively.

Run early stopping (algorithm 7.1) starting from random  $\theta$  using  $\mathbf{X}^{(\text{subtrain})}$  and  $\mathbf{y}^{(\text{subtrain})}$  for training data and  $\mathbf{X}^{(\text{valid})}$  and  $\mathbf{y}^{(\text{valid})}$  for validation data. This returns  $i^*$ , the optimal number of steps.

Set  $\theta$  to random values again.

Train on  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  for  $i^*$  steps.

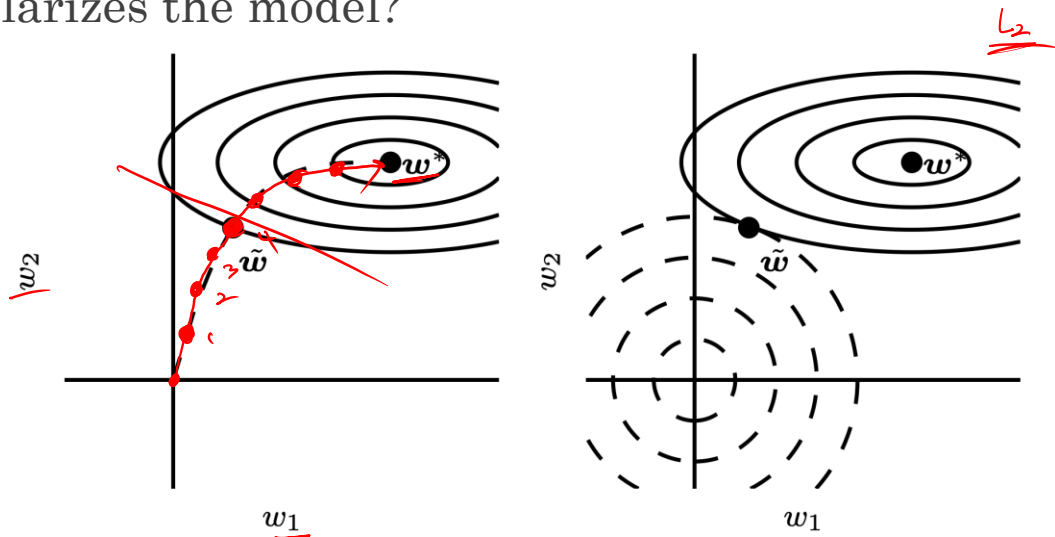
---

17

17

## Early Stopping

- What is the actual mechanism by which early stopping regularizes the model?

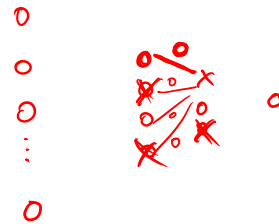


18

18

# Sparse Representations

- Meaning of  $L^1$  (or  $L^2$ ) regularization?
  - Many of the parameters become zero (or close to zero)
- $\hat{y} = f(\underline{h}) = f(f(x))$
- Representational sparsity describes a representation where many of the elements of the representation are zero (or close to zero)
- $\tilde{J}(\theta; X, y) = J + \alpha \Omega(\underline{h})$ 
  - For example,  $\Omega(\underline{h}) = \|\underline{h}\|_1$



19

19

# Ensemble Methods

- Consider for example a set of  $k$  regression models.
  - Suppose that each model makes an error  $\epsilon_i$  on each example
  - with variances  $\mathbb{E}[\epsilon_i^2] = \underline{v}$ , covariances  $\mathbb{E}[\epsilon_i \epsilon_j] = c$
  - The error made by the average prediction:  $\frac{1}{k} \sum_i \epsilon_i$
- The expected squared error
  - $\mathbb{E} \left[ \left( \frac{1}{k} \sum_i \epsilon_i \right)^2 \right] = \frac{1}{k^2} \mathbb{E} \left[ \left( \sum_i \epsilon_i \right)^2 \right] = \frac{1}{k^2} \mathbb{E} \left[ \sum_i \epsilon_i^2 + \sum_{i \neq j} 2 \epsilon_i \epsilon_j \right]$
  - $= \frac{1}{k^2} \left( \sum_i \mathbb{E}[\epsilon_i^2] + \sum_i \sum_{j \neq i} \mathbb{E}[\epsilon_i \epsilon_j] \right) = \frac{1}{k} v + \frac{1}{k^2} k(k-1) c$

20

20

# Ensemble Methods

- The expected squared error

$$\mathbb{E} \left[ \left( \frac{1}{k} \sum_i \epsilon_i \right)^2 \right] = \frac{1}{k} v + \frac{k-1}{k} c$$

- When the errors are perfectly correlated:  $c = v$

$$\mathbb{E} \left[ \left( \frac{1}{k} \sum_i \epsilon_i \right)^2 \right] = \frac{1}{k} (2v + (k-1)v) = \underline{v}$$

- When the errors are perfectly uncorrelated:  $c = 0$

$$\mathbb{E} \left[ \left( \frac{1}{k} \sum_i \epsilon_i \right)^2 \right] = \frac{v}{k}$$

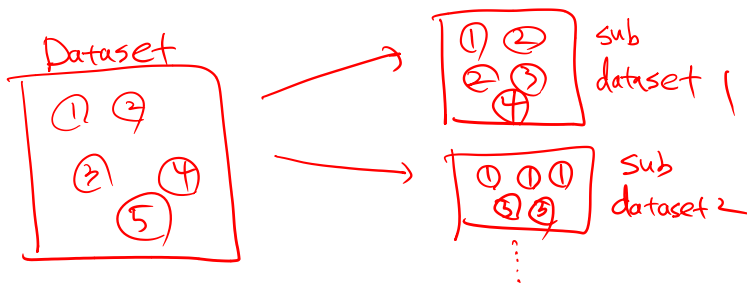
- The expected squared error of the ensemble decreases linearly with the ensemble size.

21

21

## Bagging

- Bagging involves constructing  $k$  different datasets.
  - Each dataset has the same number of examples as the original dataset
  - but each dataset is constructed by sampling with replacement from the original dataset

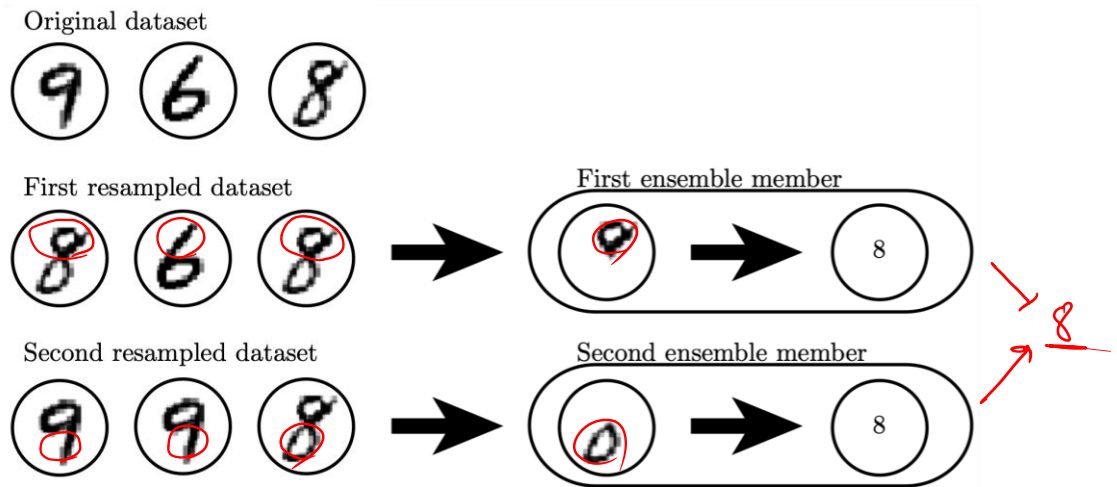


22

22

# Bagging

## ▪How bagging works



23

23

## Ensemble in Neural networks

- Neural networks reach a wide enough variety of solution points that they can often benefit from model averaging even if all of the models are trained on the **same dataset**.
  - differences in random initialization
  - random selection of minibatches,
  - differences in hyperparameters

24

24

# Dropout

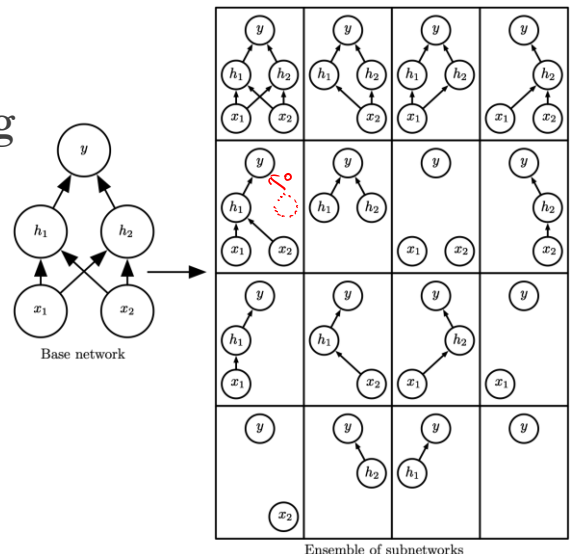
- Bagging involves training multiple models, and evaluating multiple models on each test example.
  - impractical when each model is a large neural network, since training and evaluating such networks is costly in terms of runtime and memory
- Dropout provides an inexpensive approximation to training and evaluating a bagged ensemble of exponentially many neural networks

25

25

# Dropout

- Dropout trains the ensemble consisting of all sub-networks that can be formed by removing non-output units from an underlying base network.
- In most modern neural networks, we can effectively remove a unit from a network by multiplying its output value by zero



26

26

# Dropout

---

- Learn with bagging
  - we define  $k$  different models
  - construct  $k$  different datasets by sampling from the training set with replacement
  - then train model  $i$  on dataset  $i$ .
- Dropout aims to approximate this process
  - Each time we load an example into a minibatch,
  - we randomly sample a different binary mask to apply to all of the input and hidden units in the network.
  - The mask for each unit is sampled independently.
  - The probability of sampling a mask value of one is a **hyperparameter** fixed before training begins

27

27

# Dropout

---

- The models share parameters, with each model inheriting a different subset of parameters from the parent neural network.
- The models can have different dropout probabilities for the layers.
- Predictions at test phase,
  - No dropout! (base network)
- Significant advantage of dropout is that it does not significantly limit the type of model or training procedure that can be used.

28

28