



# Convolutional Networks

---

ADVANCED ARTIFICIAL INTELLIGENCE  
JUCHEOL MOON

1

## Convolutional networks

- Convolutional networks also known as convolutional neural networks (CNN), are a specialized kind of neural network for processing data with grid-like topology.
- Time-series data, which can be thought of as a 1D grid taking samples at regular time intervals
- Image data, which can be thought of as a 2D grid of pixels.
- The name “convolutional neural network” indicates that the network employs a mathematical operation called **convolution**.

2

2

# The Convolution Operation

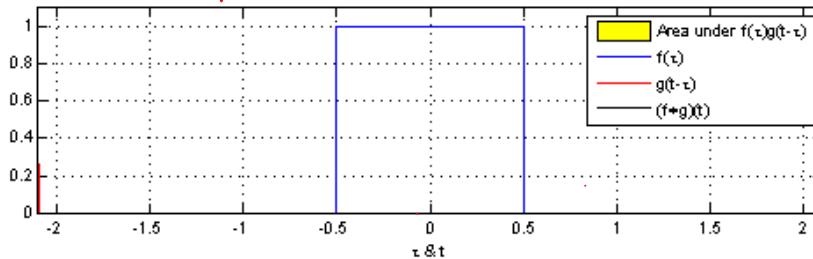
- The convolution of  $f$  and  $g$  is written  $f * g$

$$(f * g)(t) = \int_{-\infty}^{\infty} f(\tau) g(t-\tau) d\tau$$

$f$ : input

$g$ : kernel

$(f * g)$ : feature map



3

3

# The Convolution Operation

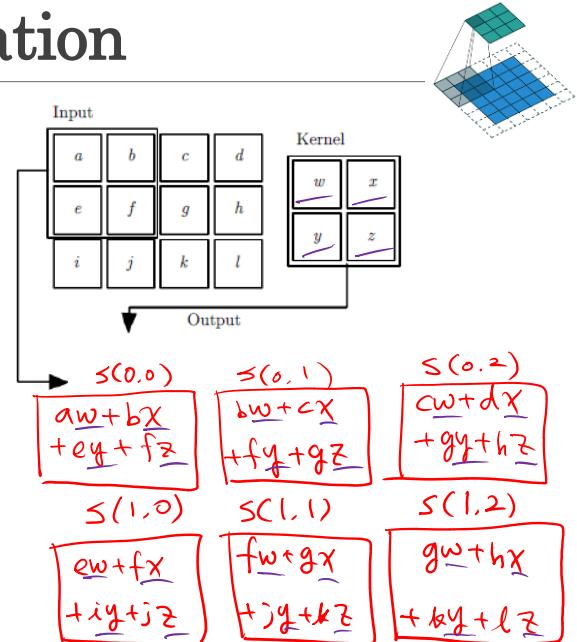
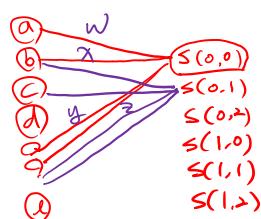
- Discretized two-dimensional

$$S(i, j) = (I * K)(i, j)$$

$$= \sum_m \sum_n I(i+m, j+n) K(m, n)$$

$I$ : input

$K$ : kernel

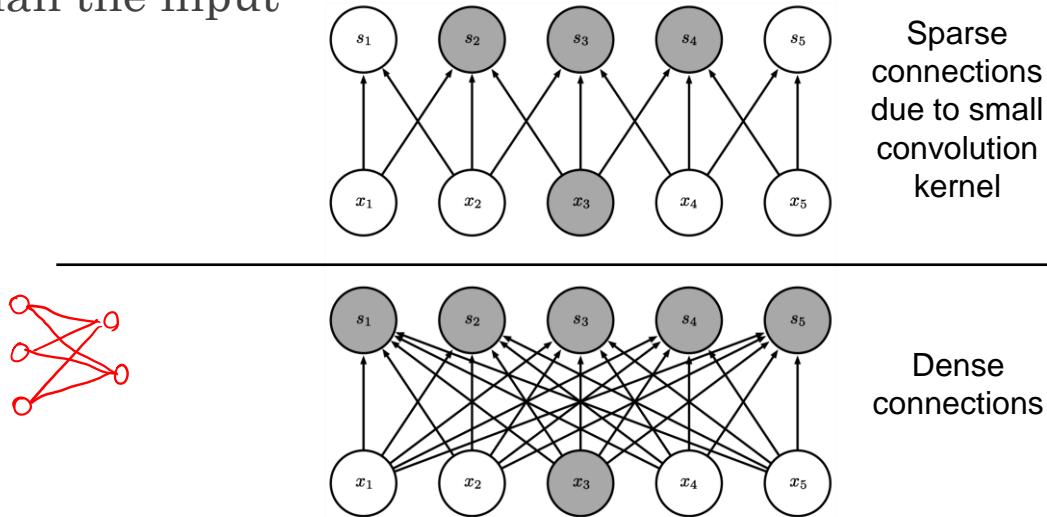


4

4

# Sparse interactions

- This is accomplished by making the kernel smaller than the input

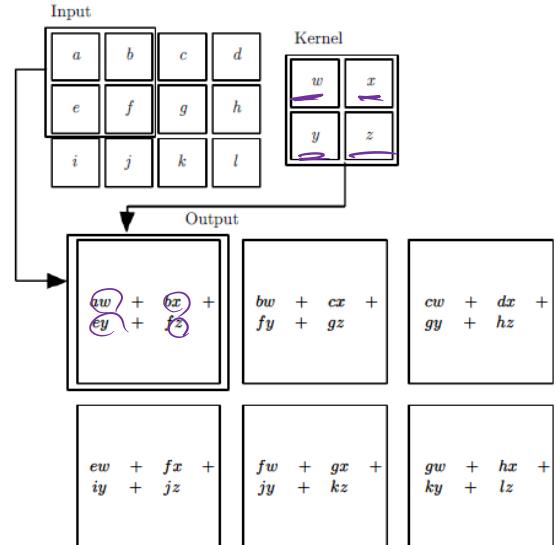


5

5

# Parameter Sharing

- Using the same parameter for more than one function in a model.
- Each member of the kernel is used at every position of the input

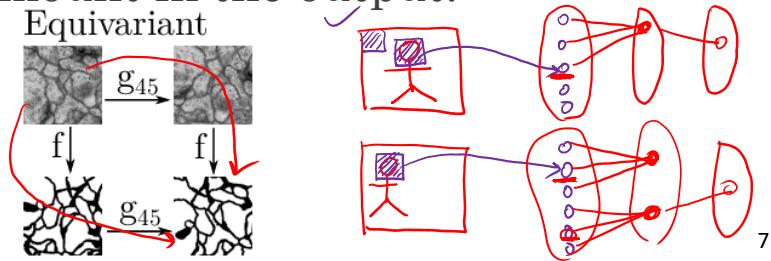


6

6

# Equivariance

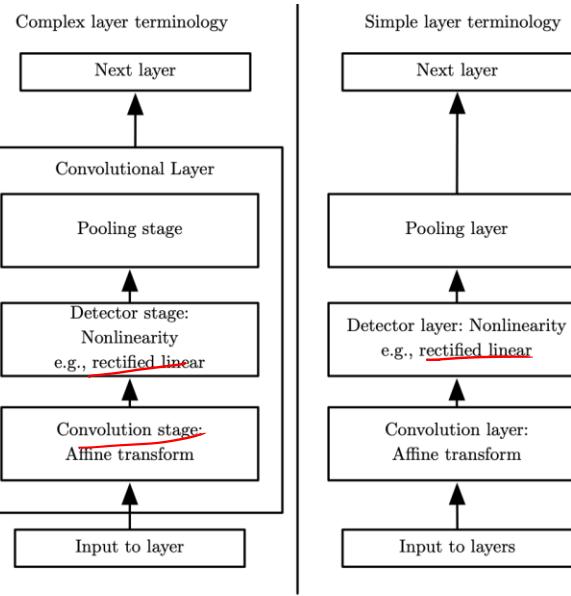
- A function  $f$  is equivariant to a function  $g$ 
  - if  $f(g(x)) = g(f(x))$
- With images, convolution creates a 2-D map of where certain features appear in the input.
- If we move the object in the input, its representation will move the same amount in the output.



7

7

# Convolutional Network Components

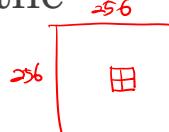


8

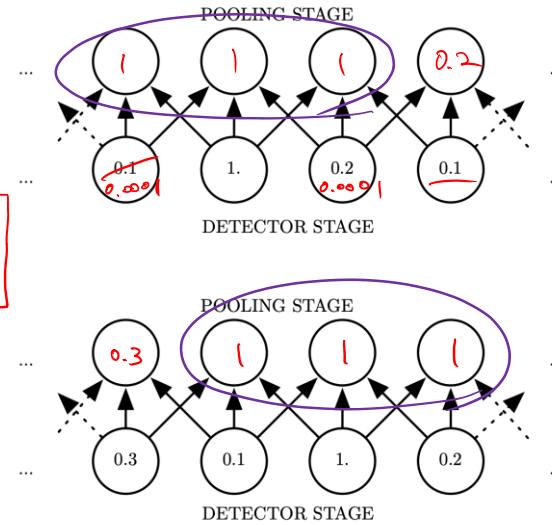
8

# Pooling

- Pooling helps to make the representation become approximately **invariant** to small translations of the input



- Invariance to local translation can be a very useful property if we care more about whether some feature is present than exactly where it is



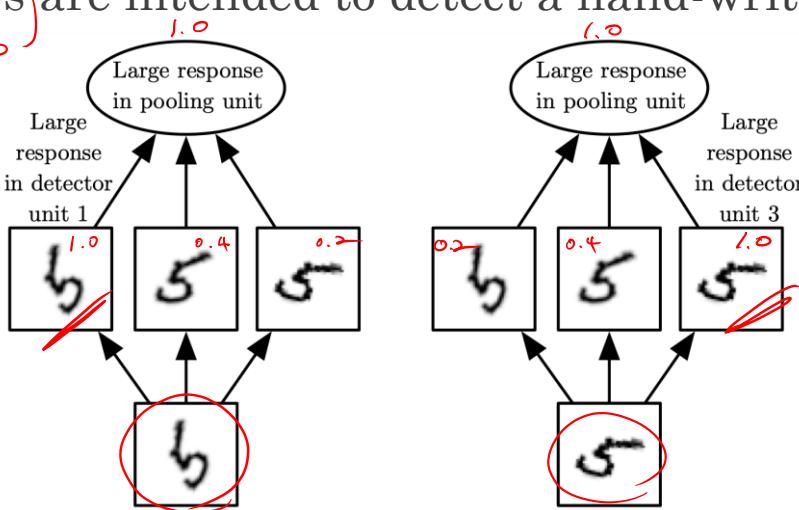
9

9

# Invariance to Learned Transformations

- Example of learned invariances
- All filters are intended to detect a hand-written 5.

(kernels)

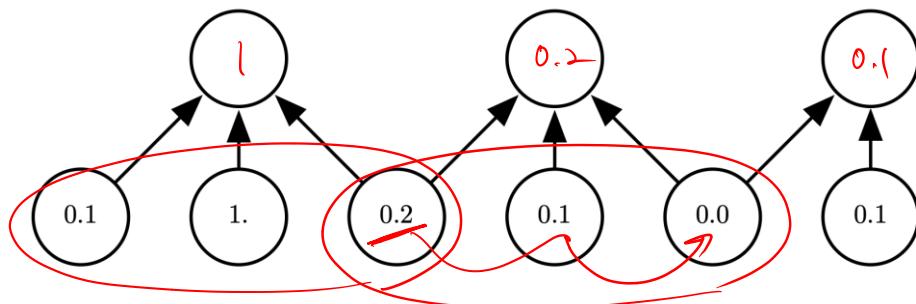


10

10

# Pooling with Downsampling

- Max-pooling with a pool width of three and a stride between pools of two
- This reduces the representation size by a factor of two, which reduces the computational and statistical burden on the next layer

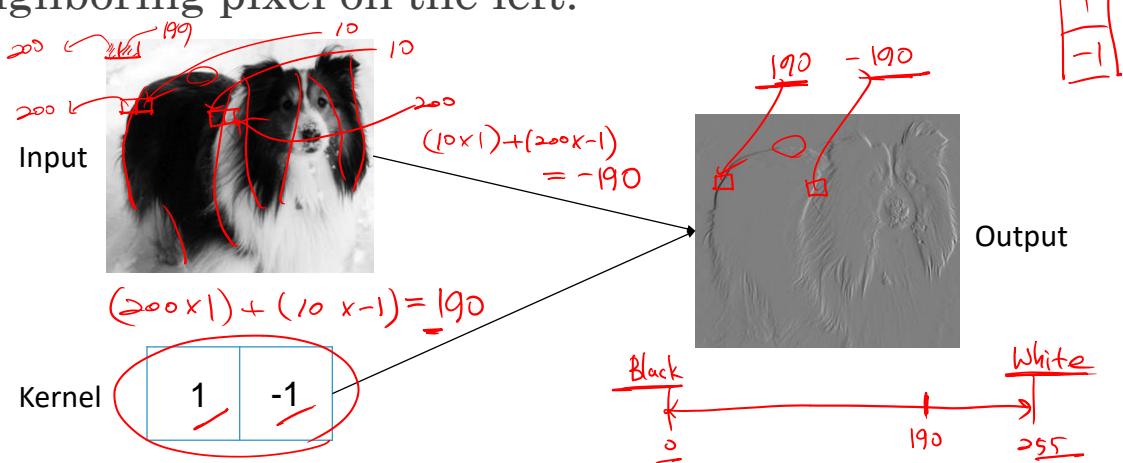


11

11

# Edge Detection by Convolution

- The image on the right was formed by taking each pixel in the original image and subtracting the value of its neighboring pixel on the left.



12

## Variants of the Basic Convolution Function

- Convolution with a single kernel can only extract one kind of feature
- The input is usually not just a grid of real values
  - A color image has a R, G, B intensity at each pixel
- When working with images, we usually think of the input and output of the convolution as being 3-D tensors, with one index into the different channels and two indices into the spatial coordinates of each channel.



13

13

## Convolution with Stride

- We may want to skip over some positions of the kernel in order to reduce the computational cost
- We can think of this as downsampling the output of the full convolution function

Input	Kernel	
a b c d	x y	$(4 \times 4) \rightarrow (3 \times 3)$
e f g h	w z	$\rightarrow \dots (1 \times 1)$
i j k l		
m n o p		

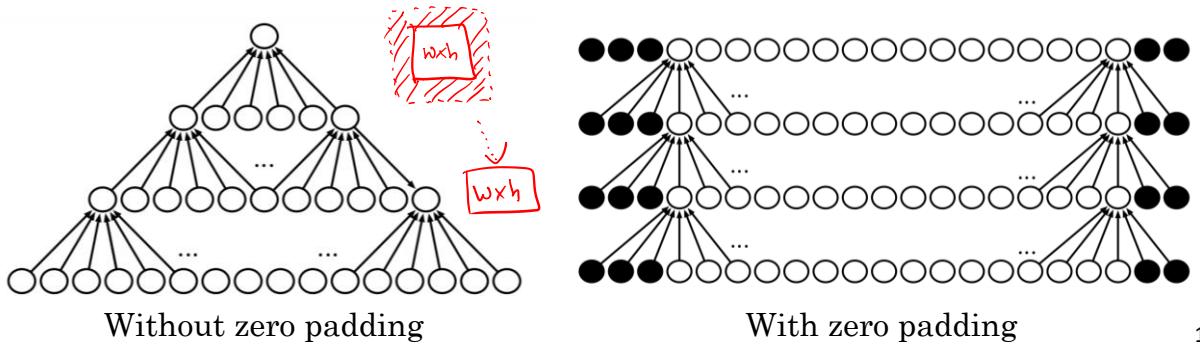
Output	
$ax+by+ew+fz$	$cx+dy+gw+hz$
$ix+jy+mw+nz$	$lx+ly+ow+pz$

14

14

# Zero Padding Controls Size

- Without zero padding, the width of the representation shrinks by one pixel less than the kernel width
- Zero padding the input allows us to control the kernel width and the size of the output independently



15

15

# Valid and Same padding in keras

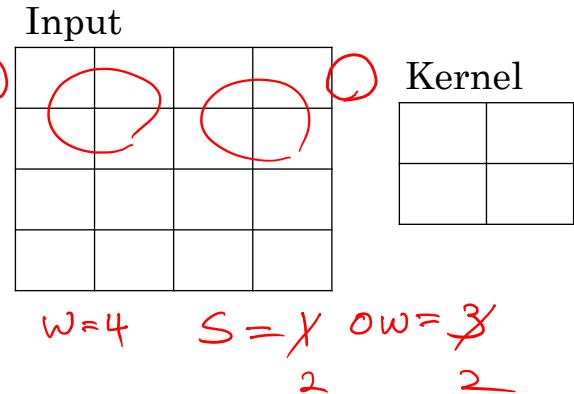
- “valid” means no padding
- The size of the output shrinks at each layer
- As layers are added, the spatial dimension of the network will eventually drop to  $1 \times 1$
- “same” results in padding the input such that the output has the same  $\text{length}$  as the original input.
- The network can contain as many convolutional layers as the available hardware can support

16

16

# Dimensions

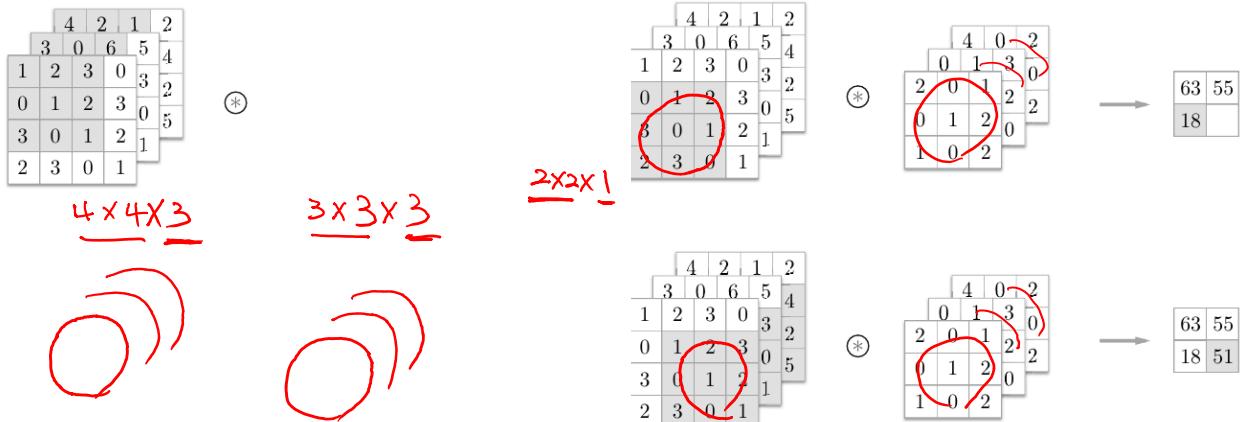
- Let padding  $P$  and stride  $S$
- The dimension of an input
  - $H \times W$
- The dimension of a kernel (filter)
  - $FH \times FW$
- The dimension of an output
  - $OH = \left\lfloor \frac{H+2P-FH}{S} \right\rfloor + 1$
  - $OW = \left\lfloor \frac{W+2P-FW}{S} \right\rfloor + 1$



17

# Convolutions on multi-channel

- 3D tensor input and kernel
  - Assume  $p = 0$  and  $s = 1$

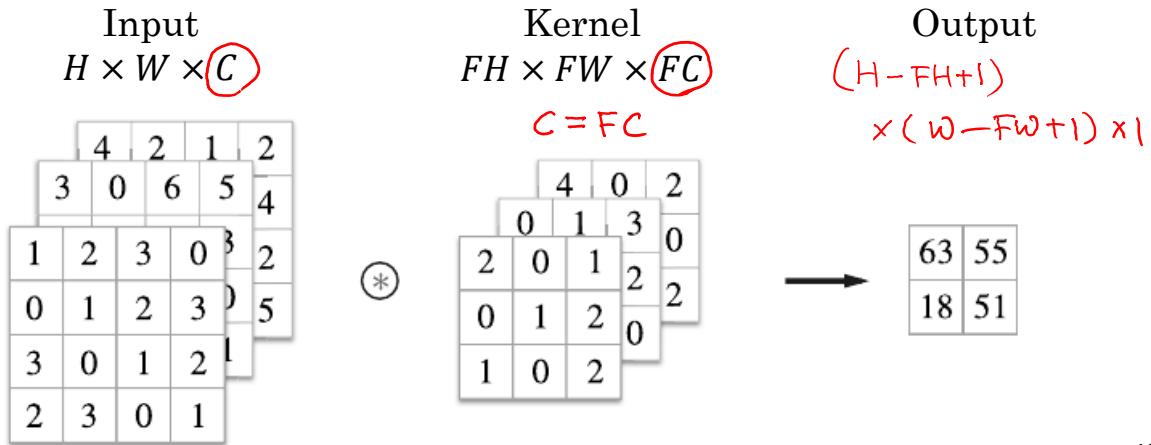


18

# Convolutions on multi-channel

- 3D tensor input and kernel
- Assume  $p = 0$  and  $s = 1$

$$OH = \left\lfloor \frac{H + p - FH}{s} \right\rfloor + 1$$

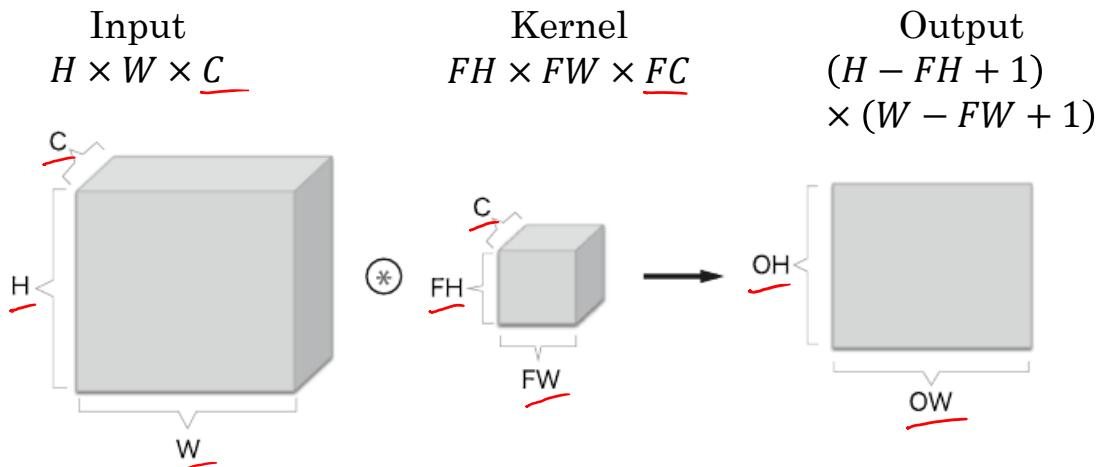


19

19

# Convolutions on multi-channel

- 3D tensor input and kernel
- Assume  $p = 0$  and  $s = 1$

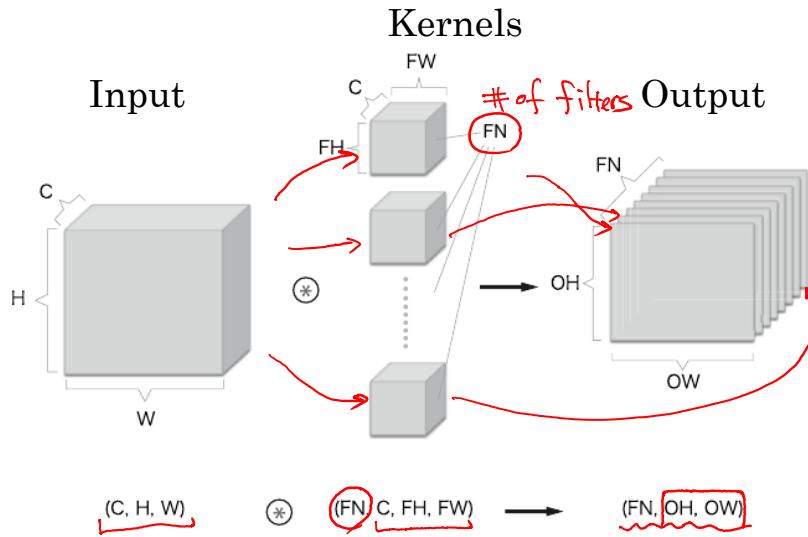


20

20

# Convolutions on multi-channel / multi-kernels

- 3D tensor input and kernels (filters)

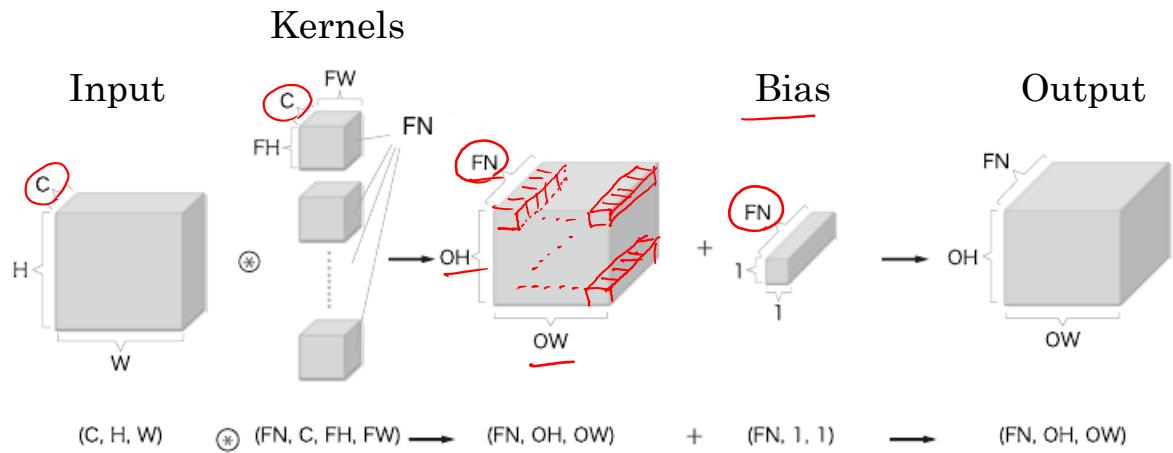


21

21

# Convolutions on multi-channel / multi-kernels

- 3D tensor input and kernel (filter)



22

22

# Max Pooling

- No parameters to be trained

1	2	1	0
0	1	2	3
3	0	1	2
2	4	0	1



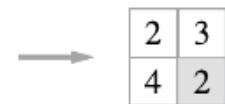
1	2	1	0
0	1	2	3
3	0	1	2
2	4	0	1



1	2	1	0
0	1	2	3
3	0	1	2
2	4	0	1



1	2	1	0
0	1	2	3
3	0	1	2
2	4	0	1



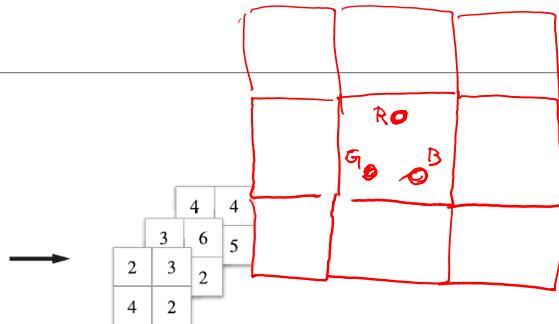
23

23

# Max Pooling

- Same number of channels

4	2	1	2
3	0	6	5
1	2	1	0
0	1	2	3
3	0	1	2
2	4	0	1



- Invariant to small translations of the input

1	2	0	7	1	0
0	9	2	3	2	3
3	0	1	2	1	2
2	3	0	1	2	1
3	2	4	0	1	0
2	6	0	1	2	1
1	2	4	0	1	8

1	1	2	0	7	1
3	0	9	2	3	2
2	3	0	1	2	1
3	2	4	0	1	0
2	6	0	1	2	1
1	2	4	0	1	8

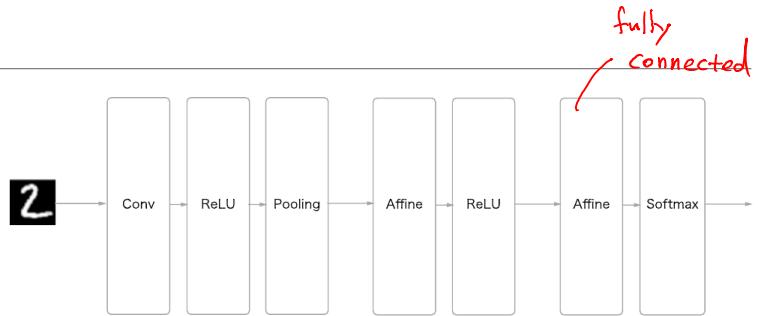
24

24

# CNN example

## ■ Design

- Input:  $28 \times 28 \times 1$
- Conv:  $5 \times 5 \times 1 \times 30$ 
  - Padding: 0
  - Stride: 1
- Pooling:  $2 \times 2 \times 1 \times 30$ 
  - Stride: 2



25

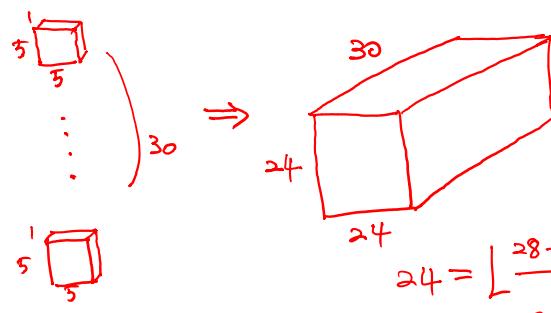
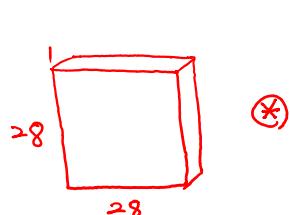
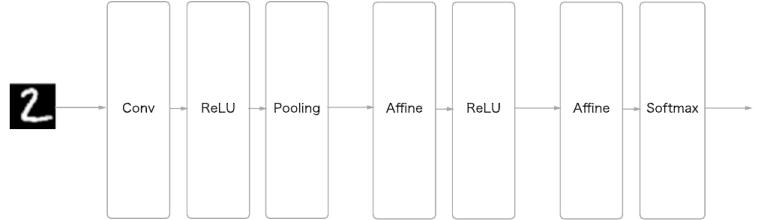
25

# CNN example

$$OH = \left\lfloor \frac{H + 2P - FH}{S} \right\rfloor + 1$$

## ■ Design

- Input:  $28 \times 28 \times 1$
- Conv:  $5 \times 5 \times 1 \times 30$ 
  - Padding: 0
  - Stride: 1



$$24 = \left\lfloor \frac{28+0-5}{1} \right\rfloor + 1$$

13

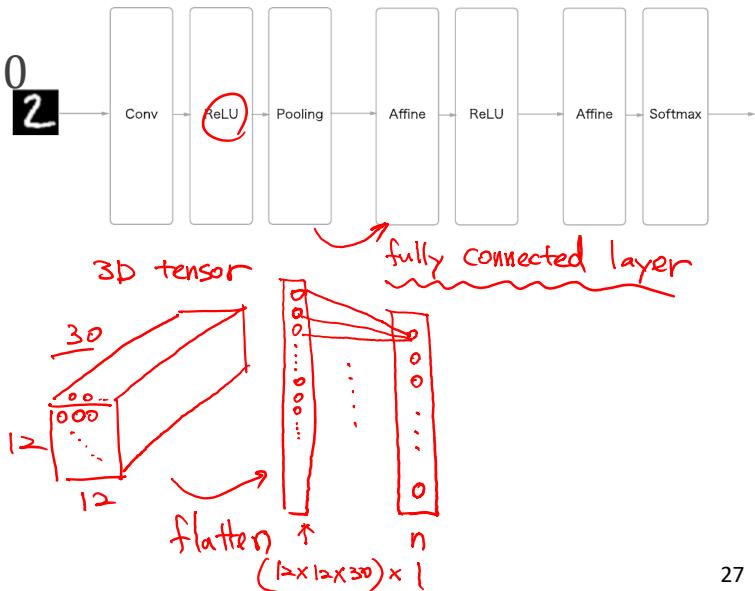
26

# CNN example

## ■ Design

- Pooling:  $2 \times 2 \times 1 \times 30$

- Stride: 2



27

27

# Bitmap

- Image pixels are generally stored with a variable number of bits per pixel which identify its color, the color depth.

- Grayscale 8 bits

- 0: Black,  $2^8 - 1 = 255$ : White



- RGB 24 bits

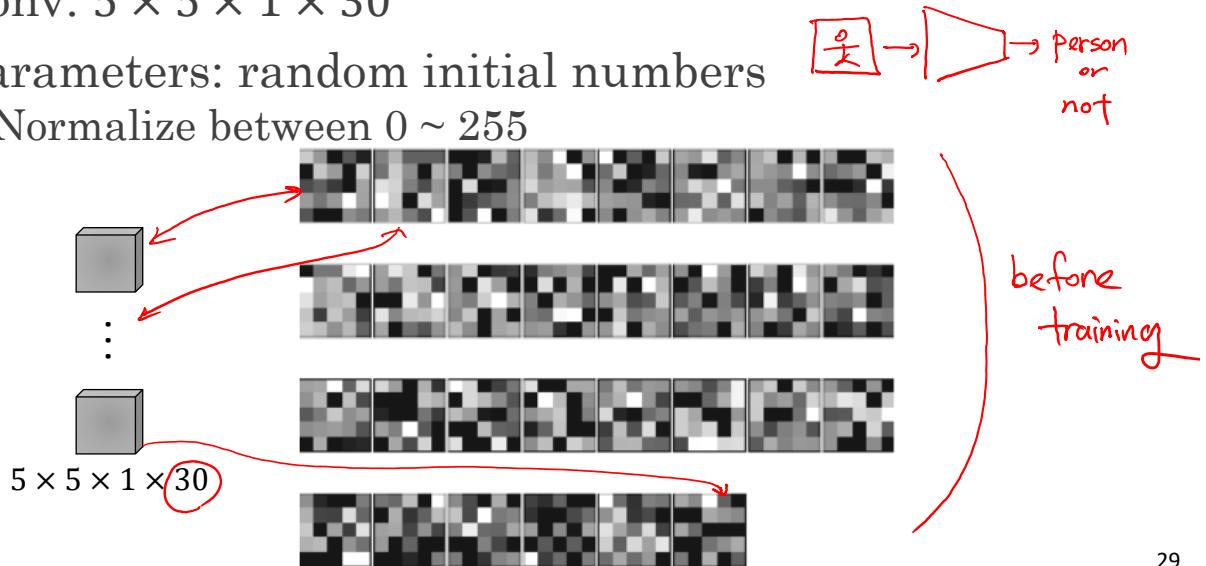


28

28

# Visualize CNN

- Conv:  $5 \times 5 \times 1 \times 30$
- Parameters: random initial numbers
  - Normalize between  $0 \sim 255$

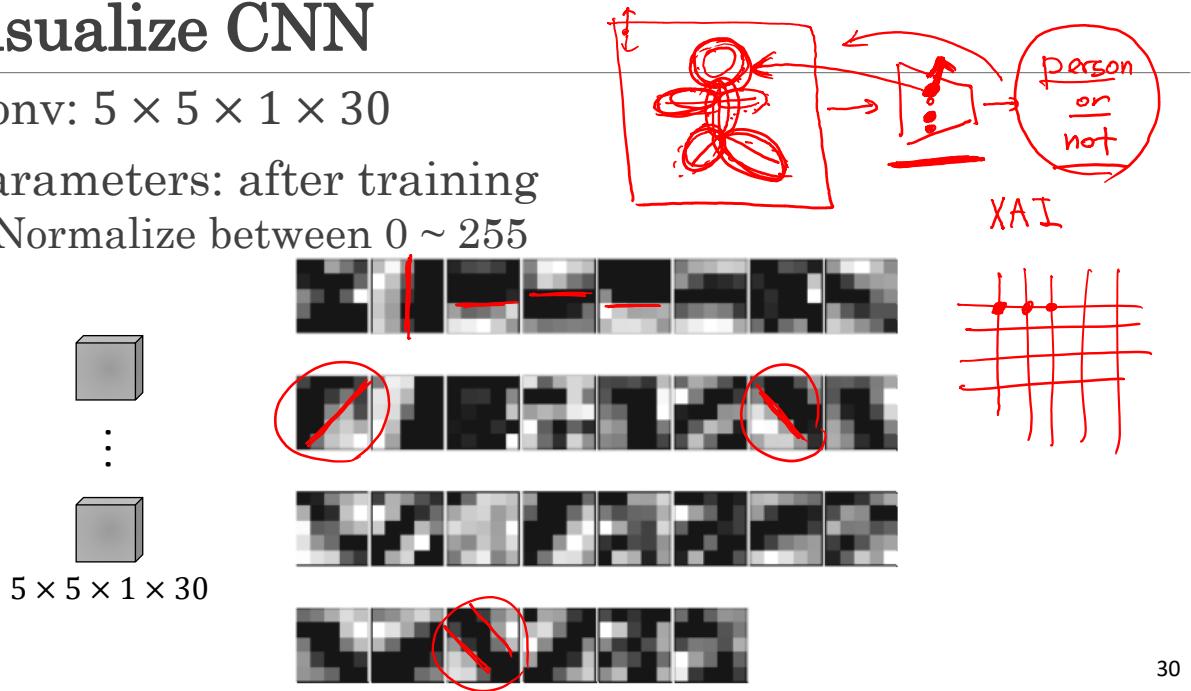


29

29

# Visualize CNN

- Conv:  $5 \times 5 \times 1 \times 30$
- Parameters: after training
  - Normalize between  $0 \sim 255$

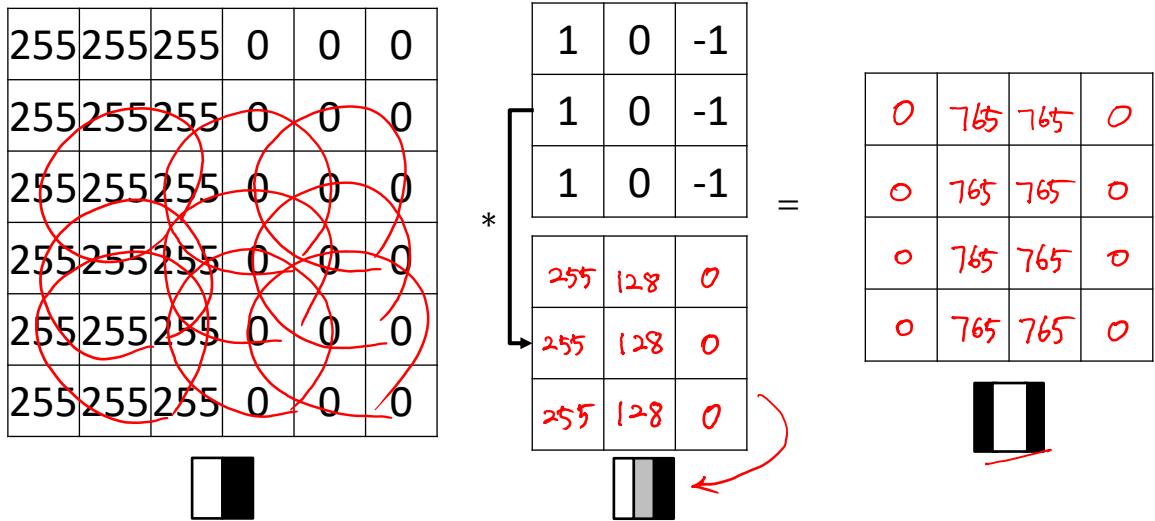


30

30

# Vertical edge detection

- How CNN detect vertical edge?

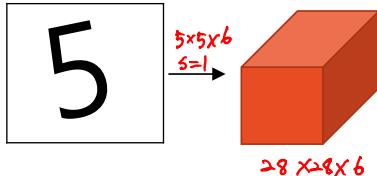


31

31

# LeNet (1998)

- Input:  $32 \times 32 \times 3$
- Conv1:  $5 \times 5 \times 3 \times 6 \rightarrow 5 \times 5 \times 6$        $OH = \left\lfloor \frac{H+2P-FH}{S} \right\rfloor + 1$
- Stride: 1, Padding: 0
- Output:  $28 \times 28 \times 6$

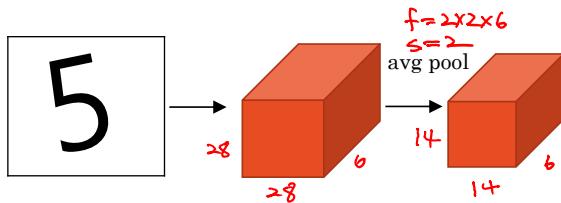


32

32

# LeNet (1998)

- Pool1:  $2 \times 2 \times 6$
- Stride: 2
- Output:  $14 \times 14 \times 6$

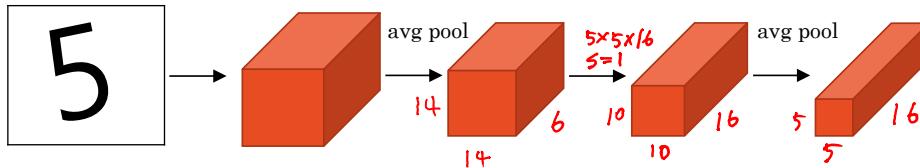


33

33

# LeNet (1998)

- Conv2:  $5 \times 5 \times 16$
- Stride: 1, Padding: 0
- Output:  $10 \times 10 \times 16$
- Pool2:  $5 \times 5 \times 16$

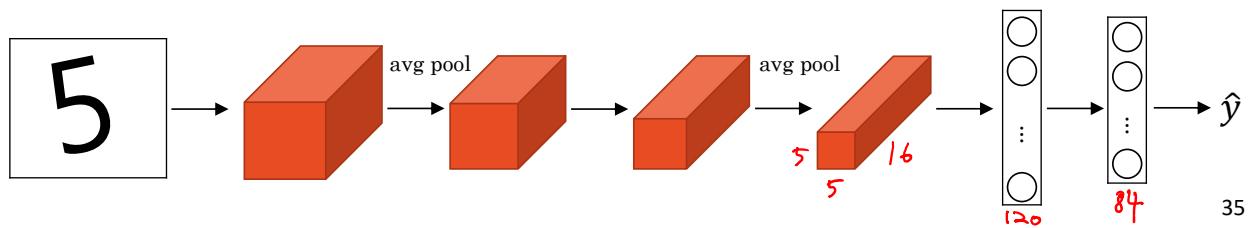


34

34

# LeNet (1998)

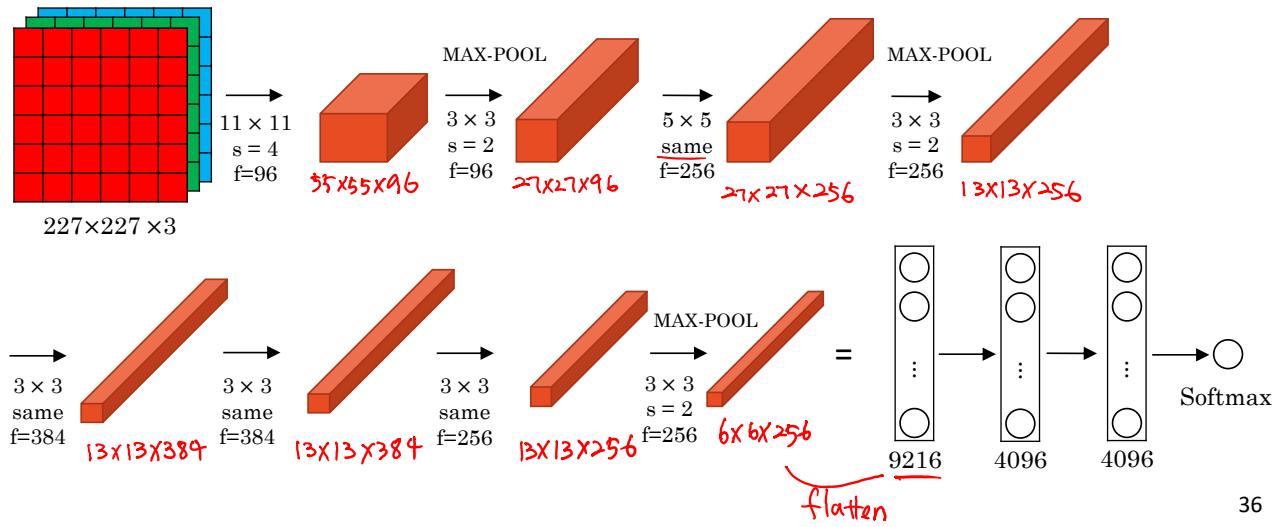
- Pool2:  $5 \times 5 \times 16$
- H3: 120
- H4: 84
- Softmax



35

# AlexNet (2012)

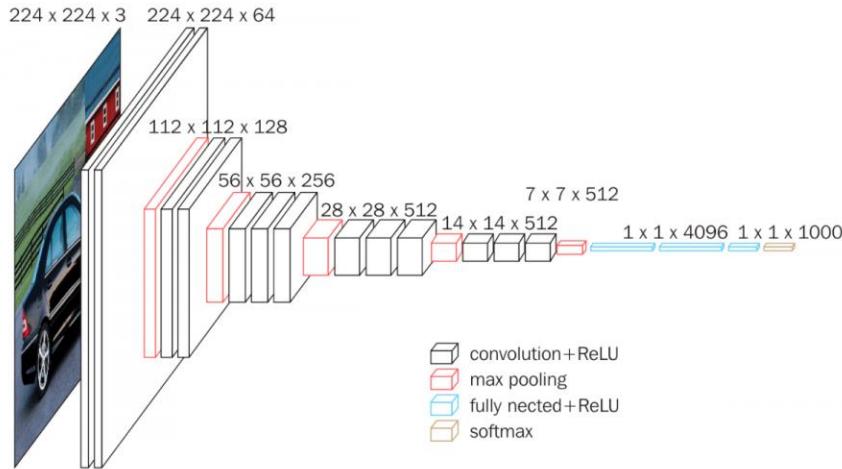
- Upgrade version of LeNet



36

# VGG-16 (2015)

```
model.add(Conv2D(filters=64,kernel_size=(3,3),padding="same", activation="relu"))
model.add(Conv2D(filters=64,kernel_size=(3,3),padding="same", activation="relu"))
model.add(MaxPool2D(pool_size=(2,2),strides=(2,2)))
```

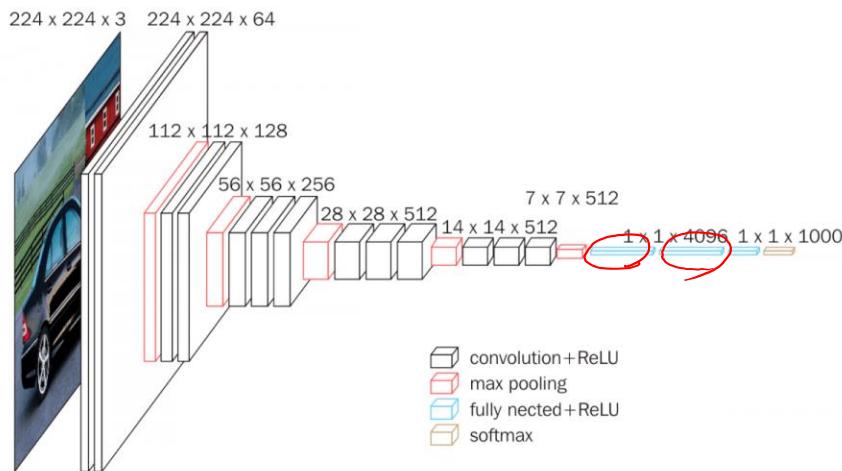


37

37

# VGG-16 (2015)

```
model.add(Flatten())
model.add(Dense(units=4096,activation="relu"))
model.add(Dense(units=4096,activation="relu"))
```

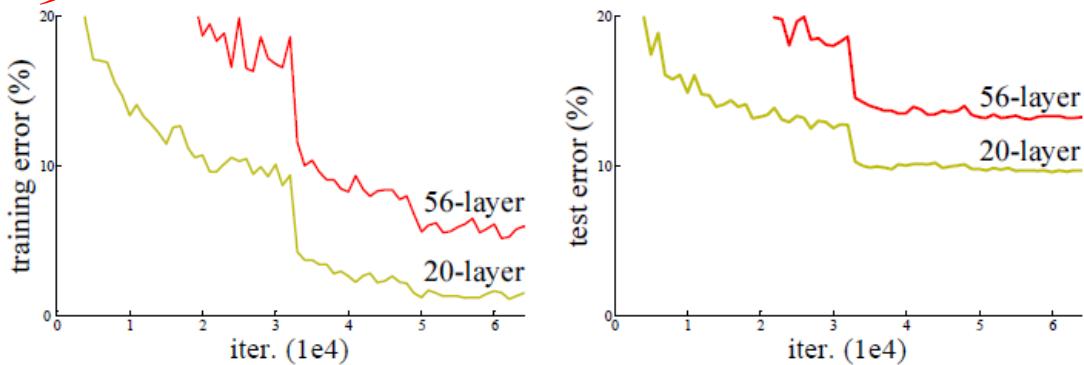


38

38

# Deep Residual Learning

- Deeper neural networks are better in general?
- Adding more layers to a suitably deep model leads to **(higher/lower)** training and test error



43

43

# Deep Neural Networks

- Can a DNN learn an identity function?

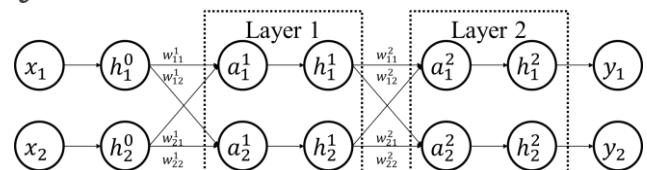
$$\vec{y} = \vec{x}$$

- Yes, but it's not easy.

$$y_1 = \vec{x}_1 \\ = h_1^1 = g(a_1^1)$$

$$\hookrightarrow w_{11}^1 h_1^0 + w_{21}^1 h_2^0$$

$$g(a_1^1) \quad g(a_2^1) \quad w_{11}^1 h_1^0 + w_{21}^1 h_2^0 \\ \hookrightarrow \frac{w_{11}^1 h_1^0 + w_{21}^1 h_2^0}{x_1} = x_1$$



$$w_{11}^1 h_1^0 + w_{21}^1 h_2^0 \\ \frac{x_1}{x_1} = x_1$$

$$= g(w_{11}^1 g(w_{11}^1 x_1 + w_{21}^1 x_2) + w_{21}^1 g(w_{21}^1 x_1 + w_{22}^1 x_2)) = x_1$$

44

44

# Deep Residual Network

- Shortcut Connections

- $a^k = w^k h^{k-1} + b^{k-1}$

- Plain block

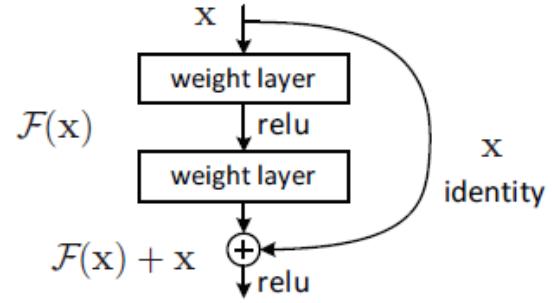
- $h^k = g(a^k)$

- Residual block

- $h^k = g(\underbrace{a^k + h^{k-1}}_{(w^k h^{k-1} + b^{k-1})} + h^k)$

$$= g((w^k h^{k-1} + b^{k-1}) + h^k)$$

$w \rightarrow 0, b \rightarrow 0$



45

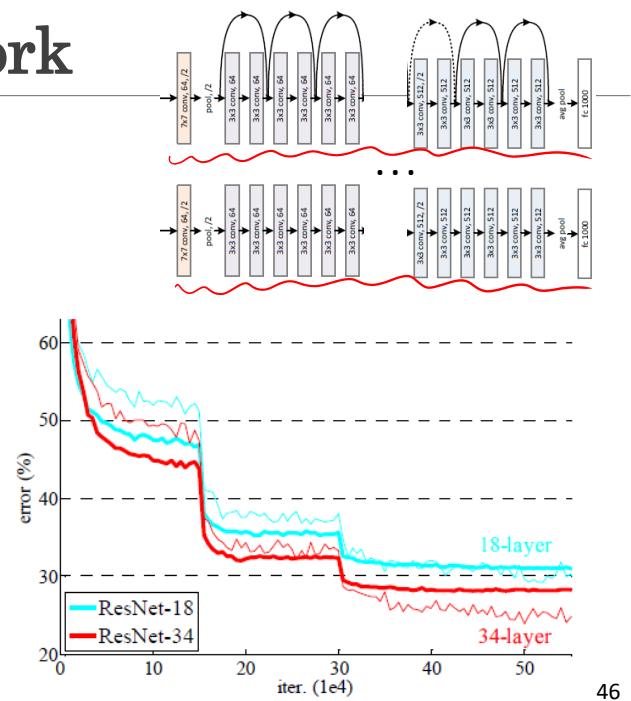
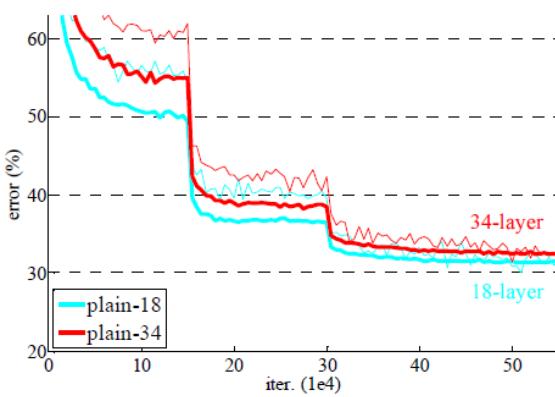
45

# Deep Residual Network

- Experiments

- Plain net of 18 and 34

- Residual net of 18 and 34



46

46

21

# ILSVRC

- The ImageNet Large Scale Visual Recognition Challenge (ILSVRC) evaluates algorithms for object detection and image classification at large scale.

Siberian husky



Eskimo dog



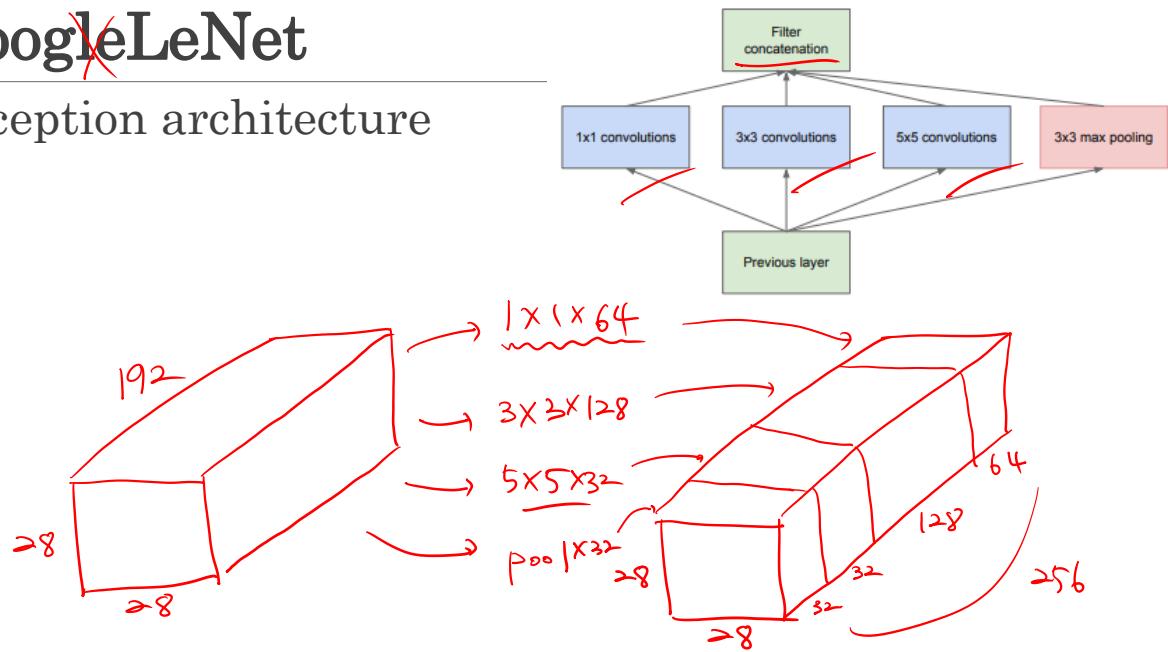
Year	Error
2012	16.4%
2013	11.7%
2014	6.67%

47

47

# GoogleLeNet

- Inception architecture



48

48

# $1 \times 1$ Convolution

- Meaning of using  $1 \times 1$  convolution

$$\begin{array}{c}
 \begin{array}{|c|c|c|} \hline 1 & 7 & 3 \\ \hline 6 & 2 & 9 \\ \hline 5 & 8 & 4 \\ \hline \end{array} & \cdot & \cdot & = & \cdot & \cdot \\
 \cdot & \cdot & \cdot & & \cdot & \cdot \\
 \end{array}$$

1

$*$

9

$\rightarrow$

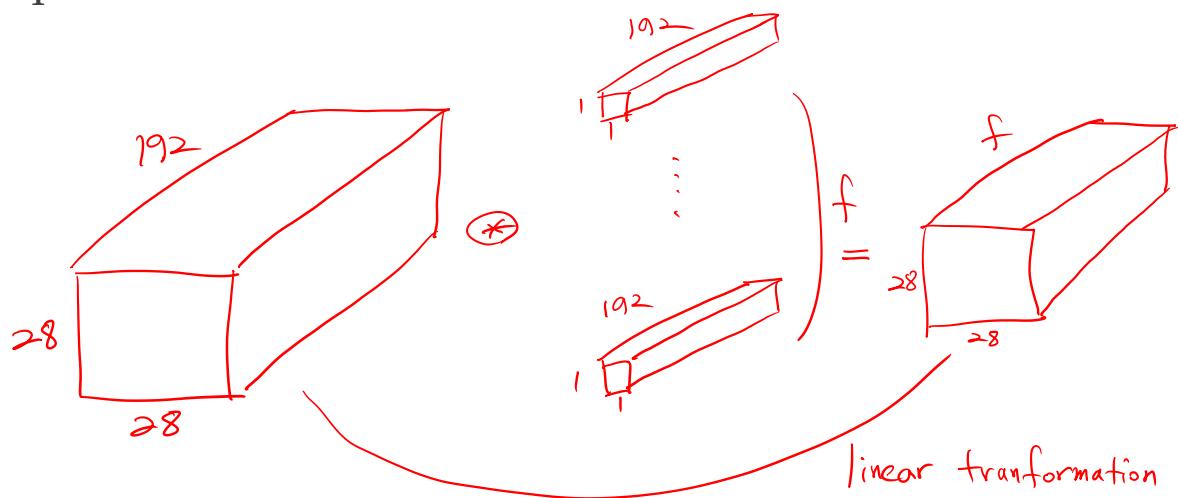
$$\begin{array}{|c|c|c|} \hline 9 & 63 & 27 \\ \hline 54 & 18 & 81 \\ \hline 45 & 72 & 36 \\ \hline \end{array}$$

49

49

# $1 \times 1$ Convolution

- Depth reduction

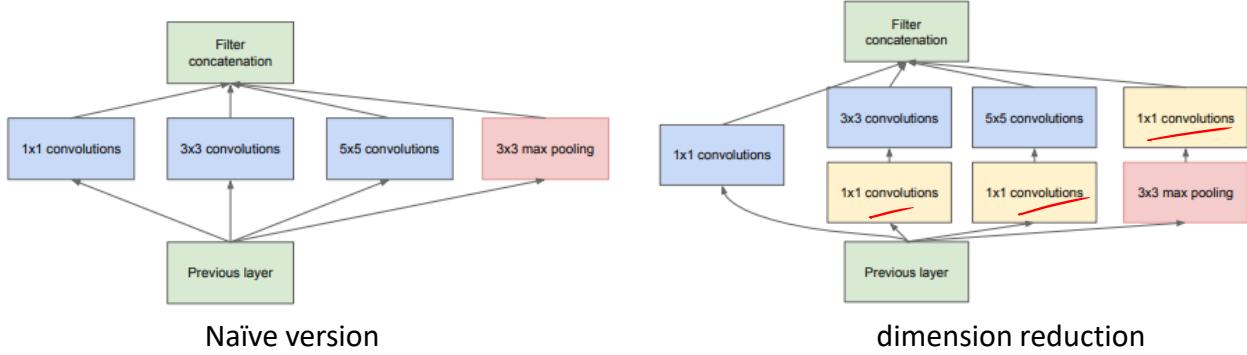


50

50

# The problem of inception network

- Computational cost

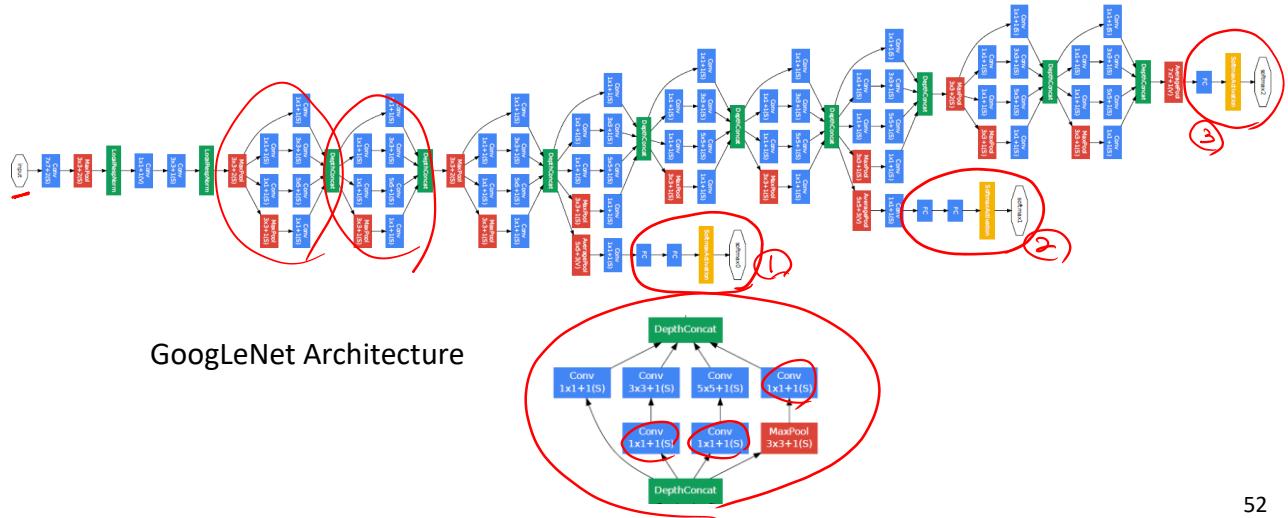


51

51

# ILSVRC 2014

- GoogLeNet

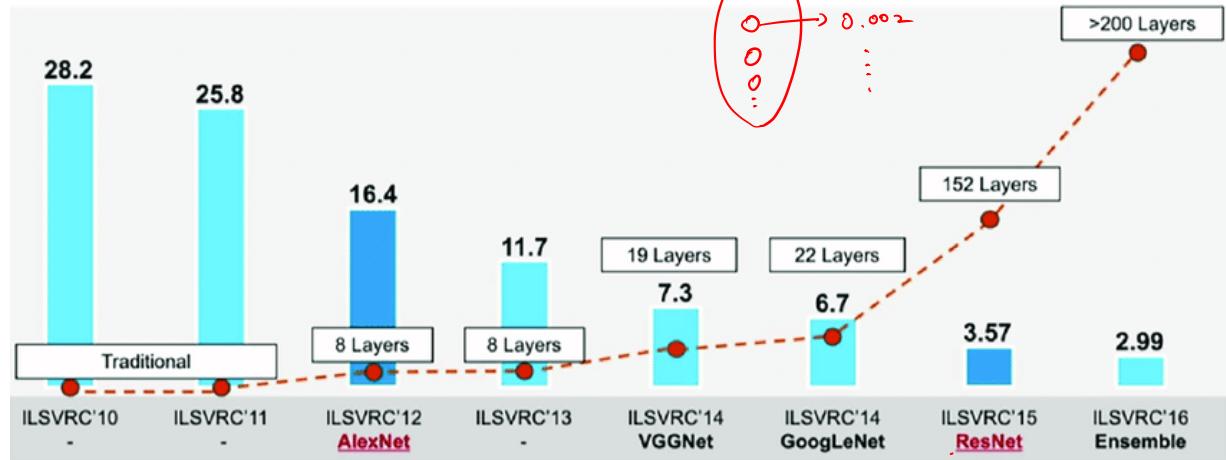


52

52

# ILSVRC

- Deeper and Deeper

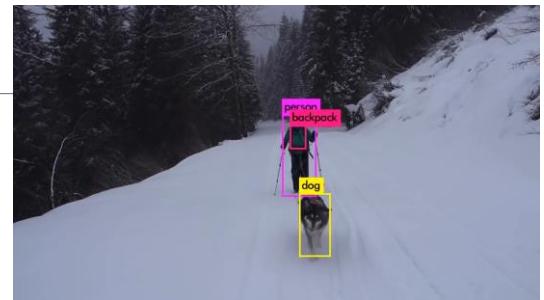


53

53

## Object Detection

- Detecting instances of semantic objects of a certain class
- Classification



*Car*

- Detection

*Car  
Location*



54

54

25

# Object Detection

- Bounding box

- $\vec{y} = [b_x, b_y, b_h, b_w, c_1, c_2, \dots, c_n]^T$

- $c_i$ : 1 if its class is  $i$

- Loss function

- $L(\vec{y}, \hat{\vec{y}}) = \|\vec{y} - \hat{\vec{y}}\|_2^2$



$$[n, n, n, n, 0.02, 0.123, \textcircled{0.134}, \dots]$$

55

55

# Object Detection

- Just when you think you have the correct answers, things get changed

- Bounding box

- $\vec{y} = [p_c, b_x, b_y, b_h, b_w, c_1, c_2, \dots, c_n]^T$

- $p_c$ : Image has the object

- $c_i$ : 1 if its class is  $i$



- Loss function

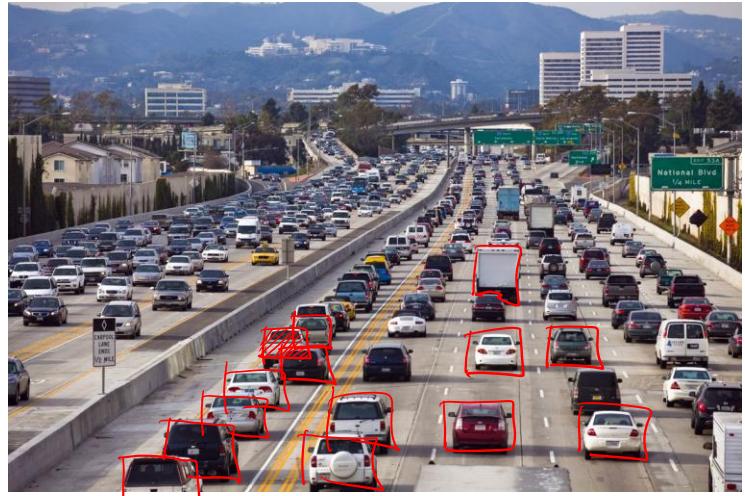
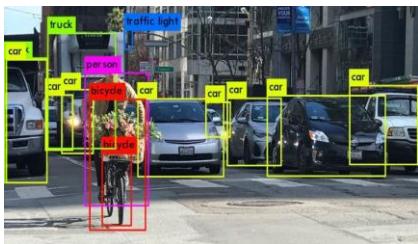
- $L(\vec{y}, \hat{\vec{y}}) = \begin{cases} \|\vec{y} - \hat{\vec{y}}\|_2^2 & \text{if } p_c = 1 \\ \frac{(y_{d1} - \hat{y}_{d1})^2}{p_c} & \text{otherwise} \end{cases}$

56

56

# Multiple Objects

- In practice, labeling job is painful.



57

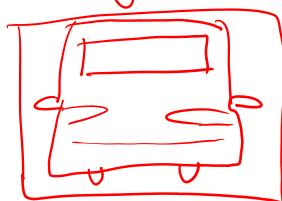
57

## Can we utilize CIFAR-10 dataset?

- We need a magic function s.t.  $f(\vec{x}) = \hat{\vec{y}}$

$\vec{x}$ : *Image*

$\hat{\vec{y}}$ : *Class label + bounding box*



Test Image



Training Images

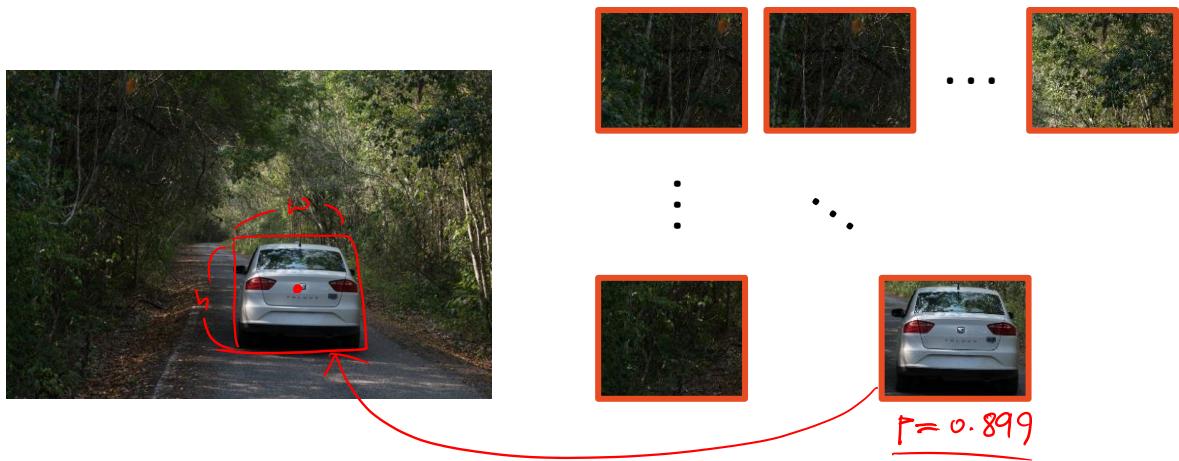


58

58

# Sliding window

- Slide a box image and classify each image inside a box



59

59

# Sliding window

- Questions?
  - Shape of the boxes
  - Size of the boxes
  - Step interval

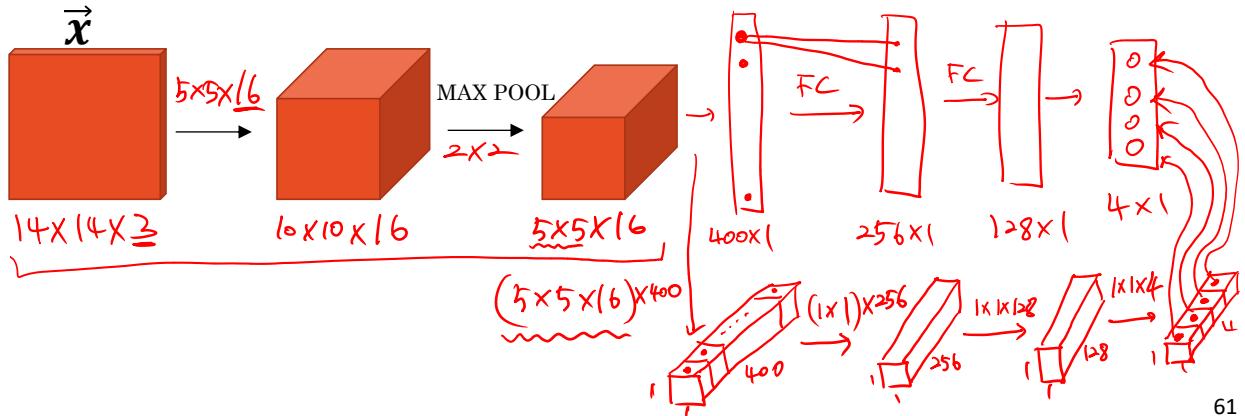


60

60

# Fully connected layers

- Typical CNN architectures include
  - Convolutional layers (Pooling layers)
  - Fully connected layers (Softmax layer)



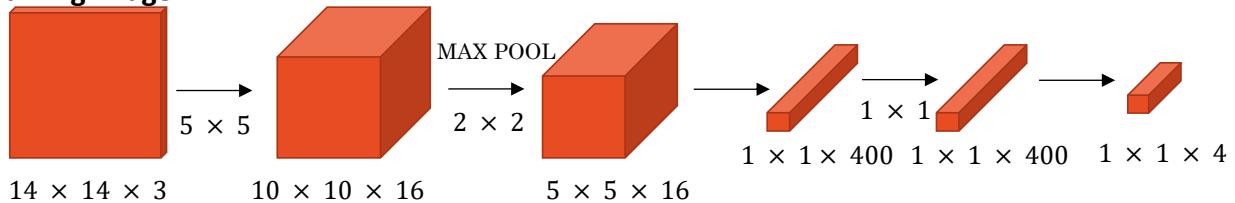
61

61

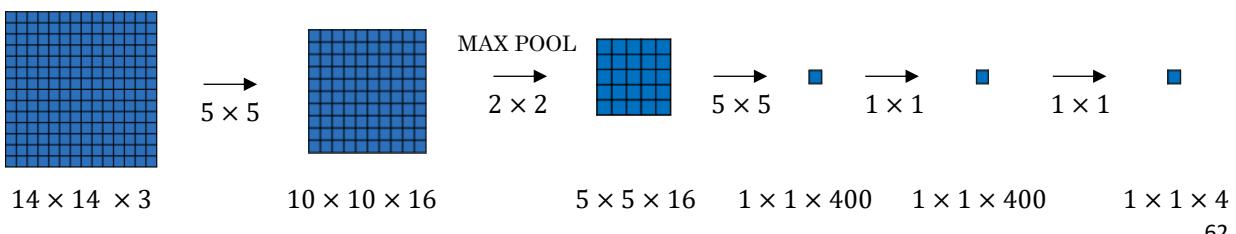
# ConvNets for detection

- Fast/cheap implementation of sliding window

**Training image**



**Simplified diagram**



62

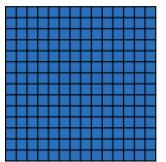
62

# ConvNets for detection

$$OH = \left\lfloor \frac{W+2P-FW}{S} \right\rfloor + 1$$

- Fast/cheap implementation of sliding window

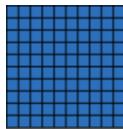
**Training image**



14 × 14 × 3

$P=0, S=1$

$\xrightarrow{5 \times 5}$



10 × 10 × 16

MAX POOL  
 $2 \times 2$



5 × 5 × 16

$\xrightarrow{5 \times 5}$

$\xrightarrow{1 \times 1}$

$\xrightarrow{1 \times 1}$

$\xrightarrow{1 \times 1}$

$\xrightarrow{1 \times 1}$

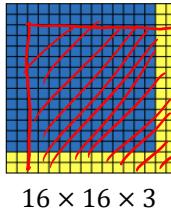
1 × 1 × 400

$\xrightarrow{1 \times 1}$

1 × 1 × 400

1 × 1 × 4

**Test image**



16 × 16 × 3

$\xrightarrow{5 \times 5}$



12 × 12 × 16

MAX POOL  
 $2 \times 2$



6 × 6 × 16

$\xrightarrow{5 \times 5}$

$\xrightarrow{1 \times 1}$

$\xrightarrow{2 \times 2 \times 400}$

$\xrightarrow{2 \times 2 \times 400}$

$\xrightarrow{2 \times 2 \times 4}$

Sermanet, et al. (2013). Overfeat: Integrated recognition, localization and detection using convolutional networks

63

63

# ConvNets for detection

- How we can change the window size and step interval?

**Training image**



10 × 10 × 3

$\xrightarrow{5 \times 5 \times 16}$

6 × 6 × 16

MAX POOL  
 $2 \times 2$



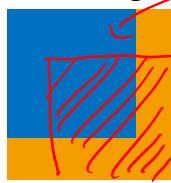
$\xrightarrow{3 \times 3 \times 400}$

$\xrightarrow{1 \times 1 \times 400}$

$\xrightarrow{1 \times 1 \times 400}$

$\xrightarrow{1 \times 1 \times 4}$

**Test image**



16 × 16 × 3

$\xrightarrow{5 \times 5}$



12 × 12 × 16

MAX POOL  
 $2 \times 2$



6 × 6 × 16

$\xrightarrow{3 \times 3}$

$\xrightarrow{1 \times 1}$

$\xrightarrow{4 \times 4 \times 400}$

$\xrightarrow{4 \times 4 \times 400}$

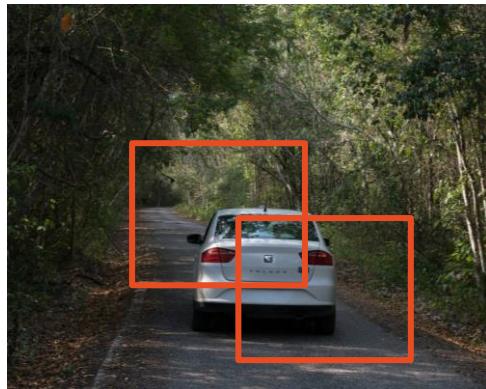
$\xrightarrow{4 \times 4 \times 4}$

64

64

## Problem of sliding window detection

- What if the bounding boxes are not matched perfectly?
  - YOLO

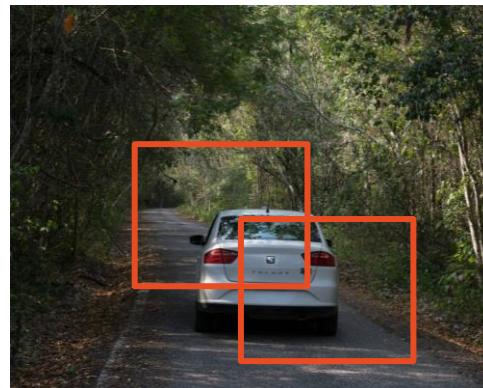


65

65

## Problem of sliding window detection

- What if the bounding boxes are not matched perfectly?
  - YOLO (You Only Live Once)

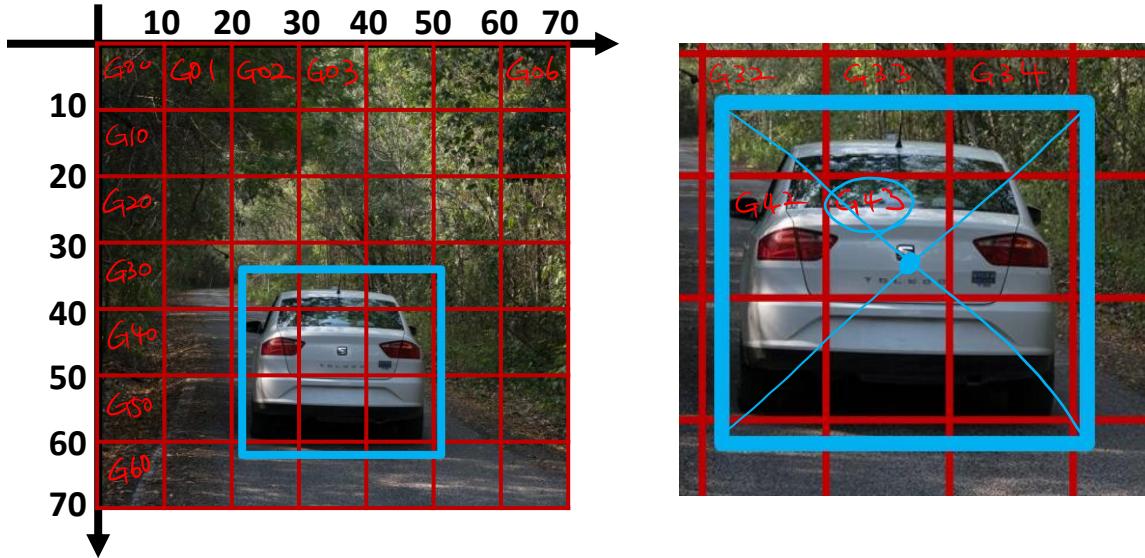


66

66

# YOLO

- To us, YOLO stands for “You Only Look Once”

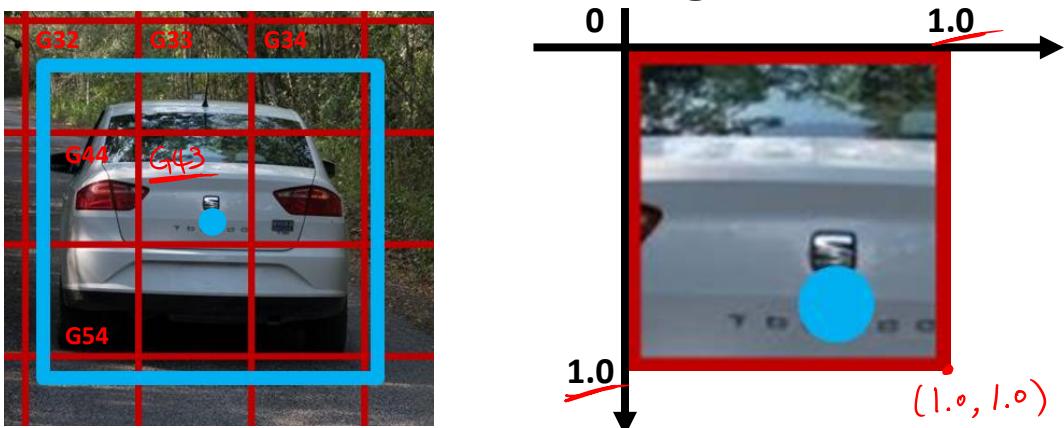


67

67

# YOLO

- Labeling for the training set
- The object is assigned to only one grid cell which includes the center of the bounding box



68

68

# YOLO

- Labeling for the training set
- The object is assigned to only one grid cell

$$\vec{y} = [P_c, b_x, b_y, b_w, b_h, c_1, c_2, \dots]^T$$

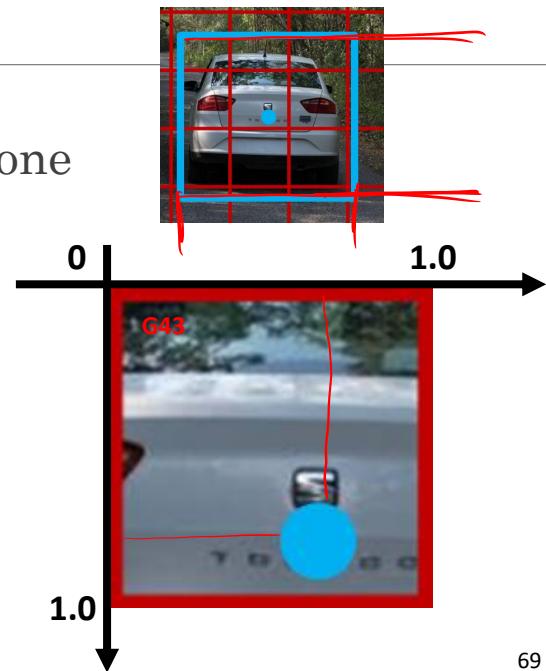
$$\vec{y}_{43} = [1, 0.9, 0.8, 2.5, 2.9, 0.1, \dots]^T$$

$$\vec{y}_{42} = [0, ?, ?, ?, \dots]^T$$

$$\vec{y}_{53} = [0, ?, ?, ?, \dots]^T$$

$$|\vec{y}| \times 7 \times 7 = |\vec{Y}|$$

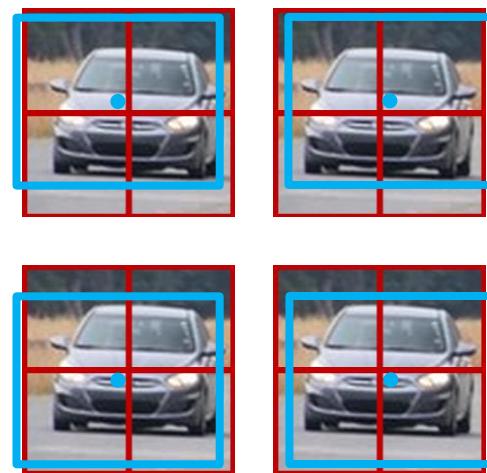
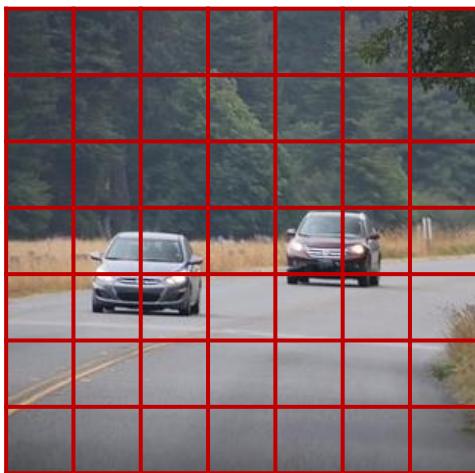
(5+4)



69

# YOLO

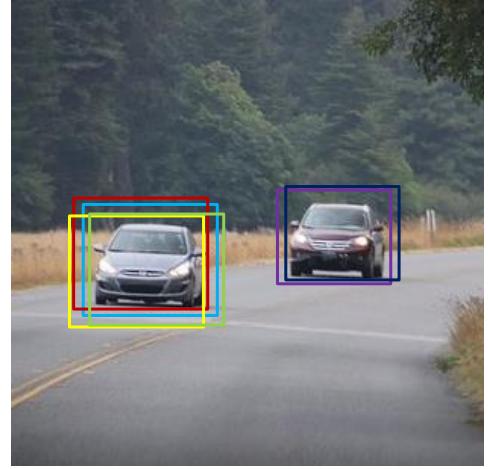
- At test phase
- Multiple grids may include the same object.



70

# YOLO

- Evaluation of bounding boxes

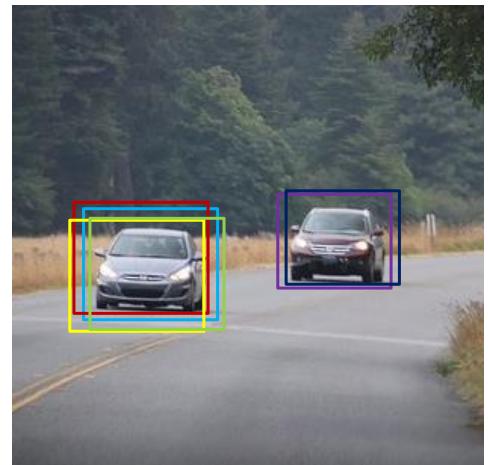


71

71

# YOLO

- Non-maximal suppression
  - Drop boxes s.t.  $p_c < \text{threshold}$
  - Group boxes that are overlapped
    - Based on intersection of union (IoU) of A and B
    - $IoU = \frac{A \cap B}{A \cup B}$
    - For instance,  $IoU > 0.5$
  - For each group, keep a box whose  $p_c$  is largest



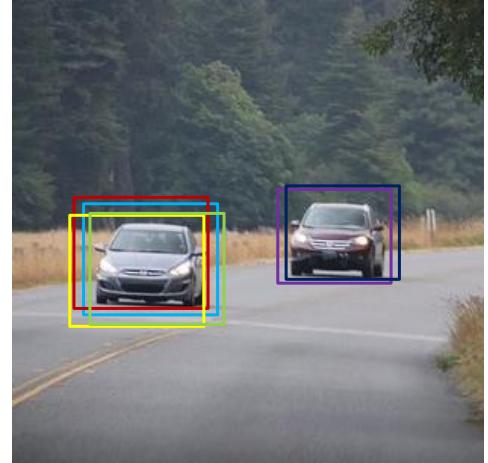
72

72

# YOLO

- Non-maximal suppression

- $\vec{y}_{red} = [0.96, \dots]^T$
- $\vec{y}_{blue} = [0.97, \dots]^T$
- $\vec{y}_{green} = [0.94, \dots]^T$
- $\vec{y}_{yellow} = [0.93, \dots]^T$
- $\vec{y}_{navy} = [0.95, \dots]^T$
- $\vec{y}_{purple} = [0.98, \dots]^T$

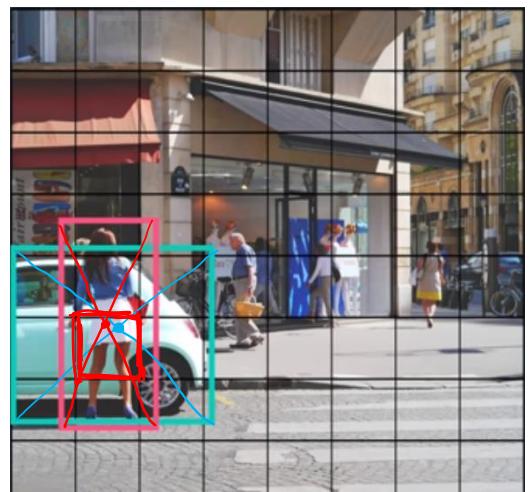


73

73

# Never enough

- It ain't over till it's over



74

74

35

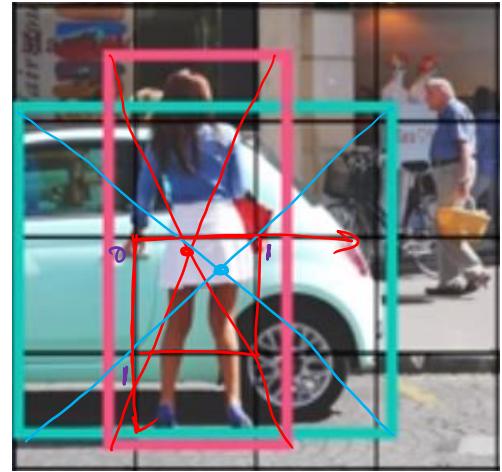
# Anchor Boxes

- A grid can have multiple objects.

$$\vec{y} = \begin{bmatrix} P_c, bx, by, bh, bw \\ P_c, bx, by, bh, bw \\ \dots \\ C_1, C_2, \dots \end{bmatrix}^T$$

ex)  $\vec{y} = \begin{bmatrix} 1, 0.4, 0.1, 3.1, 1.4 \\ 1, 0.8, 0.3, 2.8, 3.1 \\ 0, 0, \dots 1, \dots 1, 0 \end{bmatrix}$

$\uparrow$        $\uparrow$   
human    car

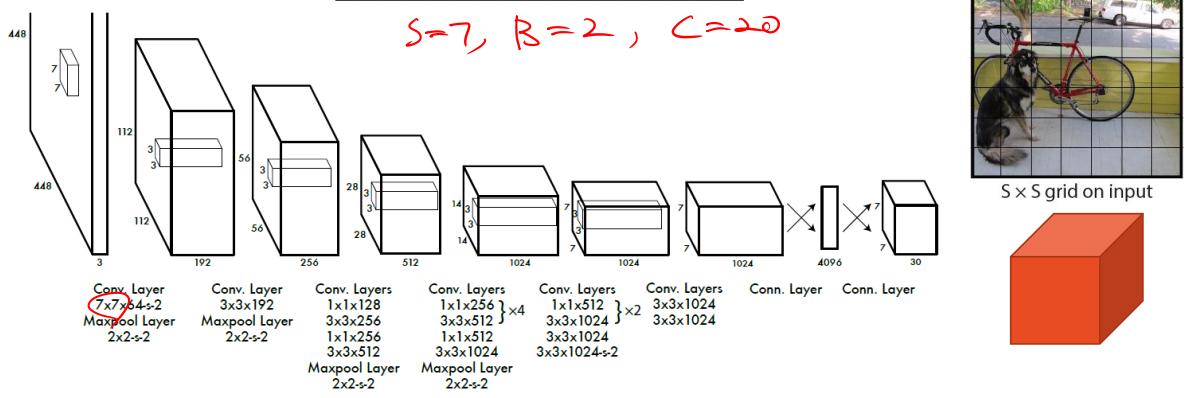


75

75

# YOLO

- YOLO divides the image into an  $S \times S$  grid and for each grid cell predicts  $B$  bounding boxes, confidence for those boxes, and  $C$  class probabilities. These predictions are encoded as an  $S \times S \times (B \times 5 + C)$  tensor.

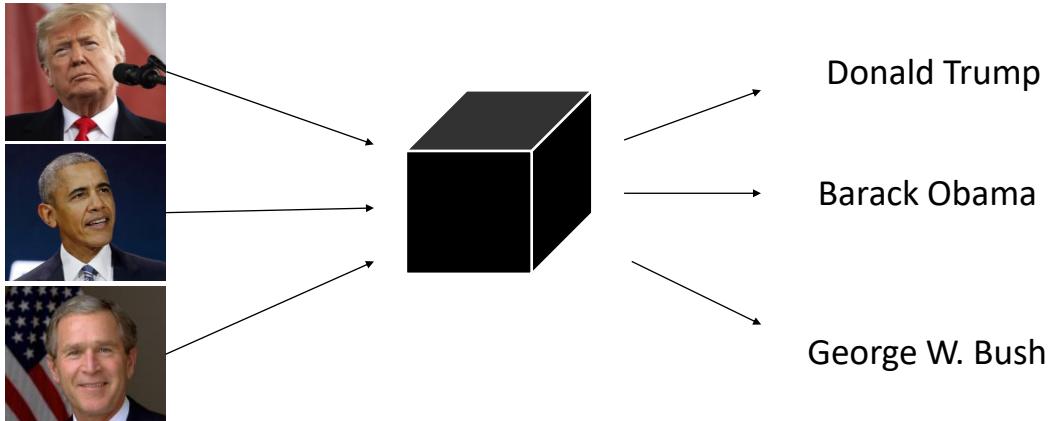


76

76

# Face Recognition

- Identifying or verifying a person from visual data.
  - Input: image / video
  - Output: id / name
- closed set recognition*

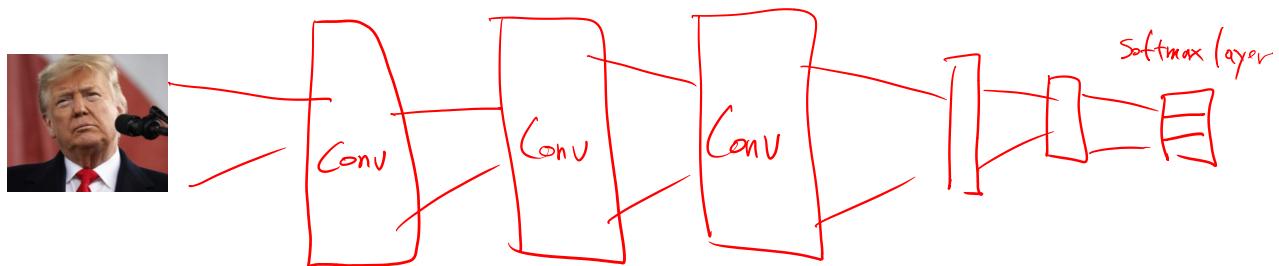


77

77

# Face Recognition

- Assume we have enough labeled data
- Network architecture (1<sup>st</sup> trial)
  - Convolutional layers (Dense layers)
  - Softmax layer

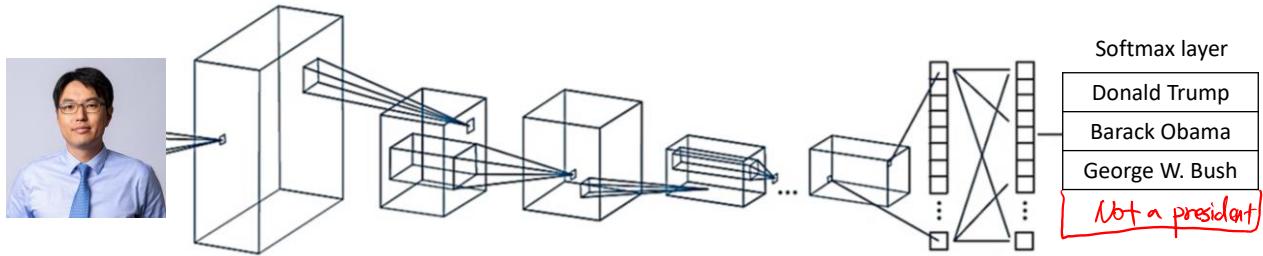


78

78

# Face Recognition

- Assume we have enough labeled data
- Network architecture (2<sup>nd</sup> trial)
  - Convolutional layers (Dense layers)
  - Softmax layer

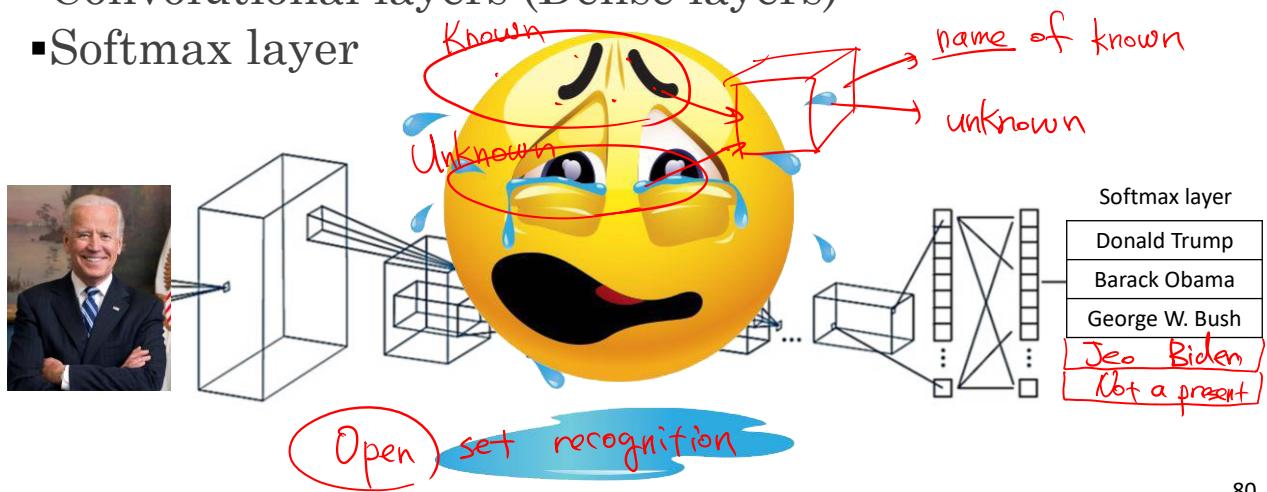


79

79

# Face Recognition

- Network architecture (3<sup>rd</sup> trial)
  - Convolutional layers (Dense layers)
  - Softmax layer



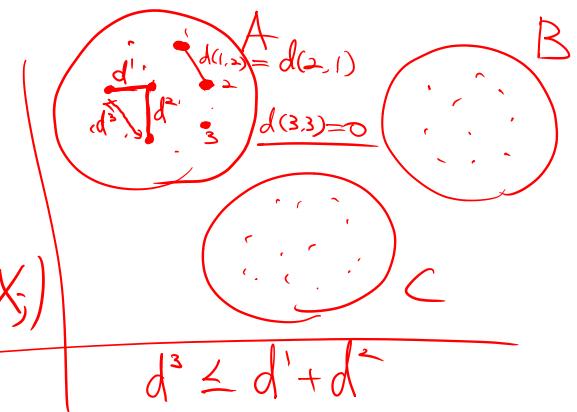
80

80

## Metric based recognition

- Let  $\vec{X}_i$  is an image of  $i$ .
- Assume we have a distance function  $d(\vec{X}_i, \vec{X}_j)$  such that

$$\left\{ \begin{array}{l} d(X_i, X_j) \geq 0 \\ d(X_i, X_j) = 0 \Leftrightarrow X_i = X_j \\ d(X_i, X_j) = d(X_j, X_i) \\ d(X_i, X_j) \leq d(X_i, X_k) + d(X_k, X_j) \end{array} \right.$$



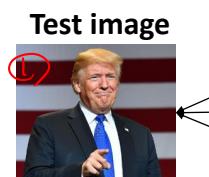
81

81

## Metric based recognition

- Assume we have a metric  $d(\vec{X}_i, \vec{X}_j)$

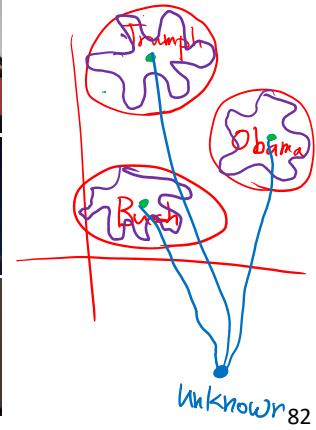
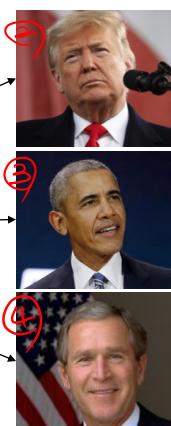
If  $d(X_i, X_j) \leq \tau \rightarrow$  recognized that person  
 else  $\Rightarrow$  unknown



Test image

$$\begin{aligned} d(1,2) & \\ d(1,3) & \\ d(1,4) & \end{aligned}$$

Training images (Known)

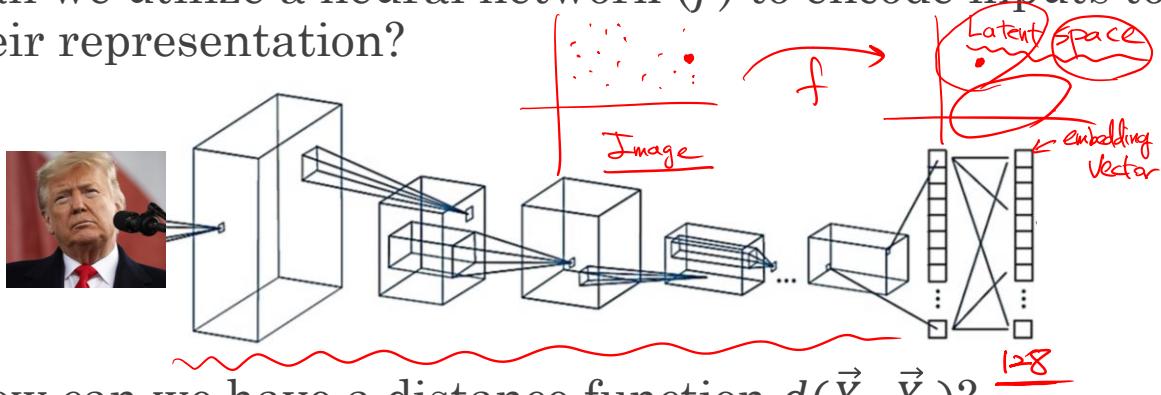


Unknown 82

82

# Representation Learning

- Can we utilize a neural network ( $f$ ) to encode inputs to their representation?



- How can we have a distance function  $d(\vec{X}_i, \vec{X}_j)$ ? 128

$$d(\vec{X}_i, \vec{X}_j) = \| f(\vec{X}_i) - f(\vec{X}_j) \|_2^2$$

$$\| \vec{X}_i - \vec{X}_j \|_2^2$$

83

83

## Embedding Vector

- We need to train a network to encode inputs such that

$$d(\vec{X}_i, \vec{X}_j) = \| f(\vec{X}_i) - f(\vec{X}_j) \|_2^2$$

is a distance between  $\vec{X}_i$  and  $\vec{X}_j$

$$\left\{ \begin{array}{l} d(X_1, X_2) < d(X_1, X_3) \\ d(X_1, X_2) < d(X_1, X_4) \\ d(X_3, X_4) < d(X_3, X_1) \\ d(X_3, X_4) < d(X_3, X_2) \\ d(X_1, X_2) ? d(X_3, X_4) \end{array} \right.$$



84

# Triplet Loss

- We want to ensure that an image  $\vec{X}^a$  (anchor) of a specific person is closer to all other images  $\vec{X}^p$  (positive) of the same person than it is to any image  $\vec{X}^n$  (negative) of any other person.

$$d(\vec{X}^a, \vec{X}^p) < d(\vec{X}^a, \vec{X}^n) \text{ for all pairs}$$

$$d(\vec{X}^a, \vec{X}^p) + \alpha < d(\vec{X}^a, \vec{X}^n) \quad \alpha \text{ is a margin.}$$

$$\underbrace{\|f(\vec{X}^a) - f(\vec{X}^p)\|_2^2}_{\text{Distance between anchor and positive}} + \alpha - \underbrace{\|f(\vec{X}^a) - f(\vec{X}^n)\|_2^2}_{\text{Distance between anchor and negative}} < 0$$

85

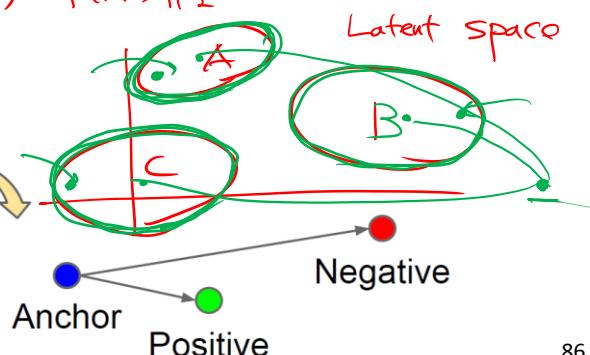
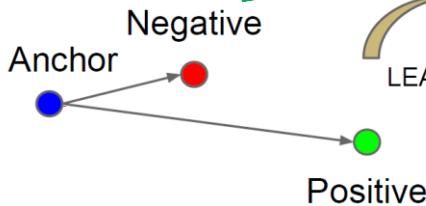
85

# Triplet Loss

- Triplet Loss minimizes the distance between an anchor and a positive and maximizes the distance between the anchor and a negative

$$L = \|f(\vec{X}^a) - f(\vec{X}^p)\|_2^2 - \|f(\vec{X}^a) - f(\vec{X}^n)\|_2^2 + \alpha$$

1. Clustering  $\xrightarrow{k\text{-means}}$
2. Centroid  $\xrightarrow{\text{OSVM}}$



86

41